

Practice using MapReduce

“Let the wild rumpus begin!”
- Maurice Sendak

ResourceManager


NameNode

MRAppMaster

DataNode

NodeManager

YarnChild



Key points from the last lecture

Name nodes know everything about the data (metadata, FS images)

- data nodes are dumb
- blocks are BIG - 128 M
- block contents are unstructured

Writing HDFS data is distributed and efficient

- pipelines are created for replication
- writes are thoroughly verified

The NameNode

- manages the “metadata” for data stored on the cluster
- monitor the health of Datanodes

The ResourceManager schedules Jobs and monitors their health

Advanced: References for HDFS

- **HDFS in more depth:**

Hadoop: The Definitive Guide, 4rd Edition, by Tom White, Chapter 3.

- **File system shell guide**

<http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

- **Namenode startup**

<http://hortonworks.com/blog/understanding-namenode-startup-operations-in-hdfs/>

- **Namenode availability: checking-pointing explained**

<http://blog.cloudera.com/blog/2014/03/a-guide-to-checkpointing-in-hadoop/>

- **HDFS Metadata directories explained**

<http://hortonworks.com/blog/hdfs-metadata-directories-explained/>

Agenda

- In-class practice 1: setup VM and work with HDFS
- What is Map Reduce
- Intro to Spark
- In-class practice 2: running Map Reduce jobs
 - Learn to run both Spark and MR2 jobs in Eclipse
 - Learn to submit Spark and MR2 jobs to a cluster

Practice 1

Setup and use the VM

Flash drive

- **Contains a file called “fall2016.ova”**
- ***Practice 1* - see install directions for fall2016.ova**
 - For the cowboys:
 - Don't unzip or try to open this file.
 - Don't try to create a “new” VM, use the [import](#) function as described in *Practice 1*.

Important: copy the .ova to your computer before you import it!

- Flashdrives have a limited number of “writes”
- If you try to run the fall2016.ova from your flashdrive it will be intolerably SLOW
- These flashdrives get **HOT. Really hot.**

The next 45 - 60 minutes ...

Import the VM

- **time to copy fall2016 to your computer: 5-10 minutes**
- **time to import fall2016.ova to your VM: 10-15 minutes**
- **time to boot VM: 5-10 minutes**

Do some things in the VM:

- **get comfortable with using the VM**
- **work with HDFS**

FOLLOW THE PRACTICE 1 instructions closely and in order...

Deep dive MapReduce

APIs and job submission

The 3 phases of a Map-Reduce job in MR2

Map phase

- **MapTask sets up the Mapper by running Mapper's**
 - **setup() method**
 - **run() method**
 - **cleanup method**
- **The Mapper's run() method executes map() for each record in a "split"**



Shuffle-sort - buckets up the mapped data, sorts and ships to the reducers

Reduce phase

- **ReduceTask merges the data buckets fetched from the Mappers**
- **ReduceTask runs Reducer's set-up, run and tear-down methods**
- **Output of each Reducer is saved as a file in HDFS**

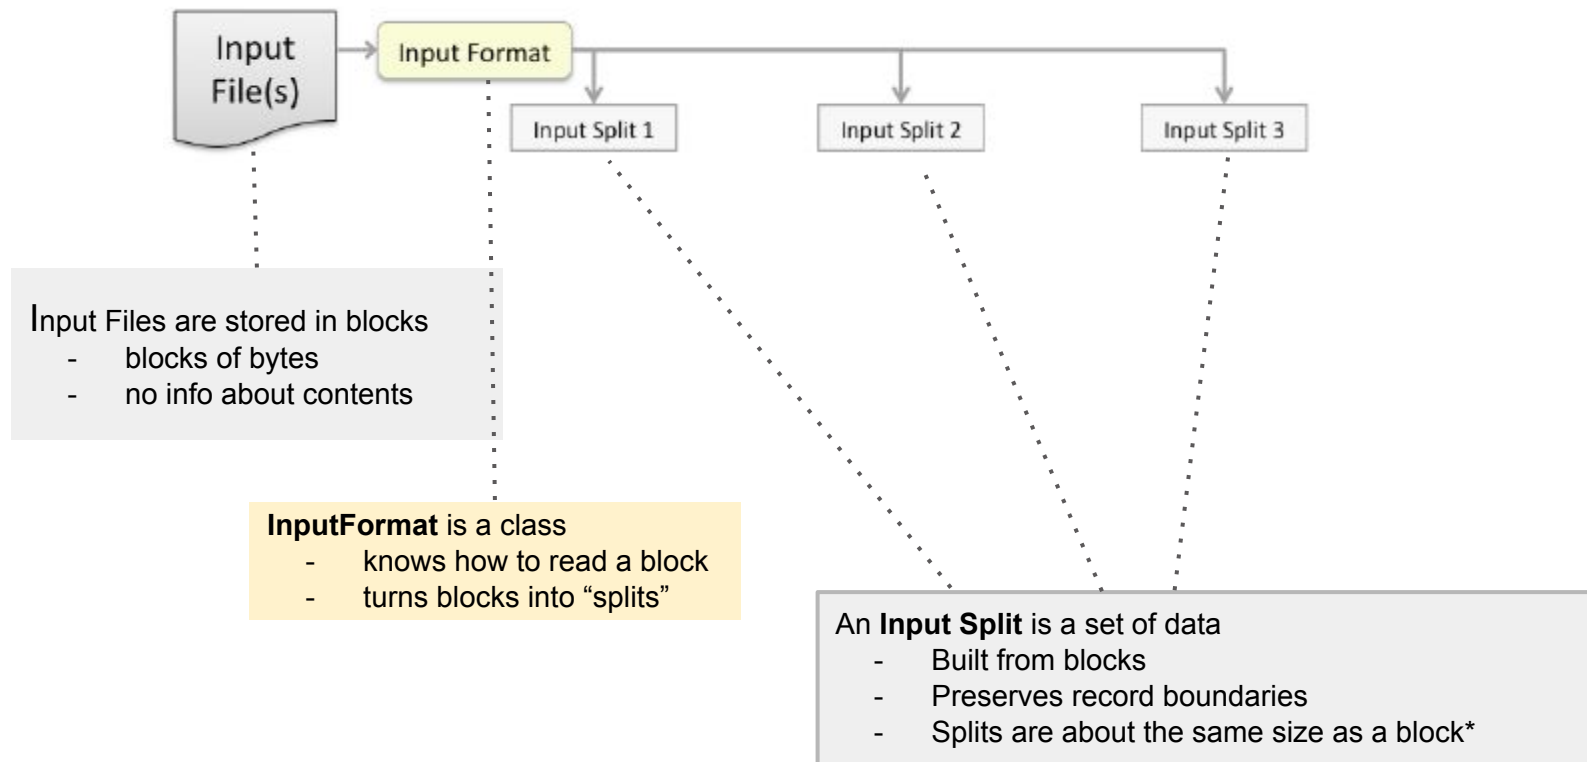
The MapReduce API for traditional MR2

Developers customize MapReduce jobs by extending:

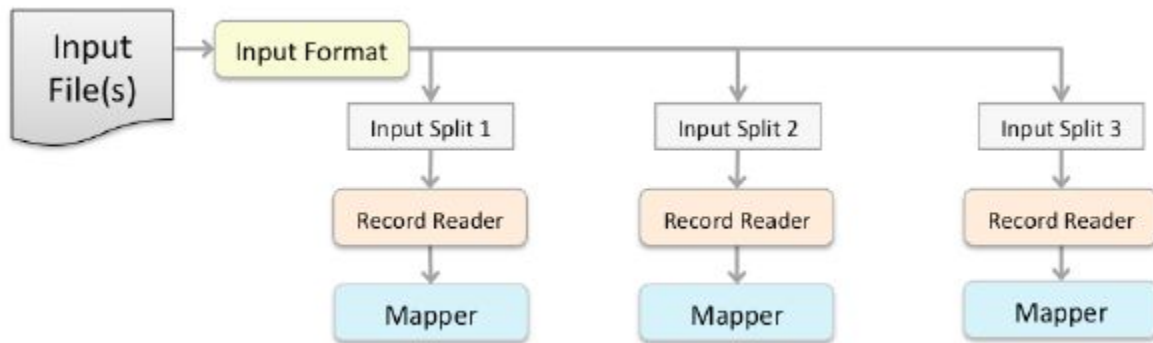
- InputFormat
 - RecordReader
 - Mapper
 - Combiner
 - Partitioner
- 
- Instantiated and run in the MapTask
- Reducer
 - OutputFormat
 - RecordWriter
- 
- Instantiated and run in the ReduceTask

These classes are exposed as part of MR2's API.

MapReduce Processing

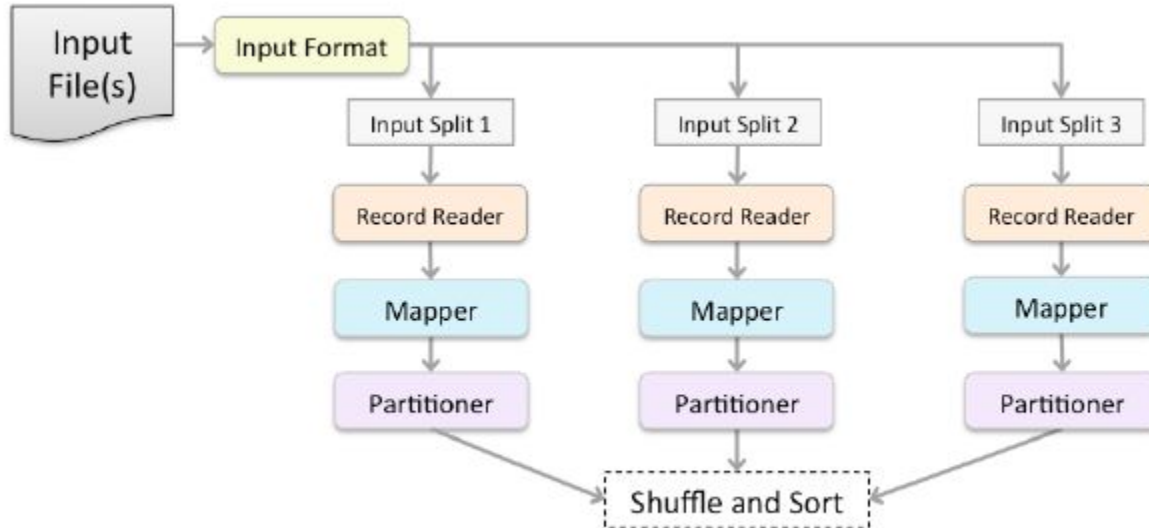


MapReduce Processing



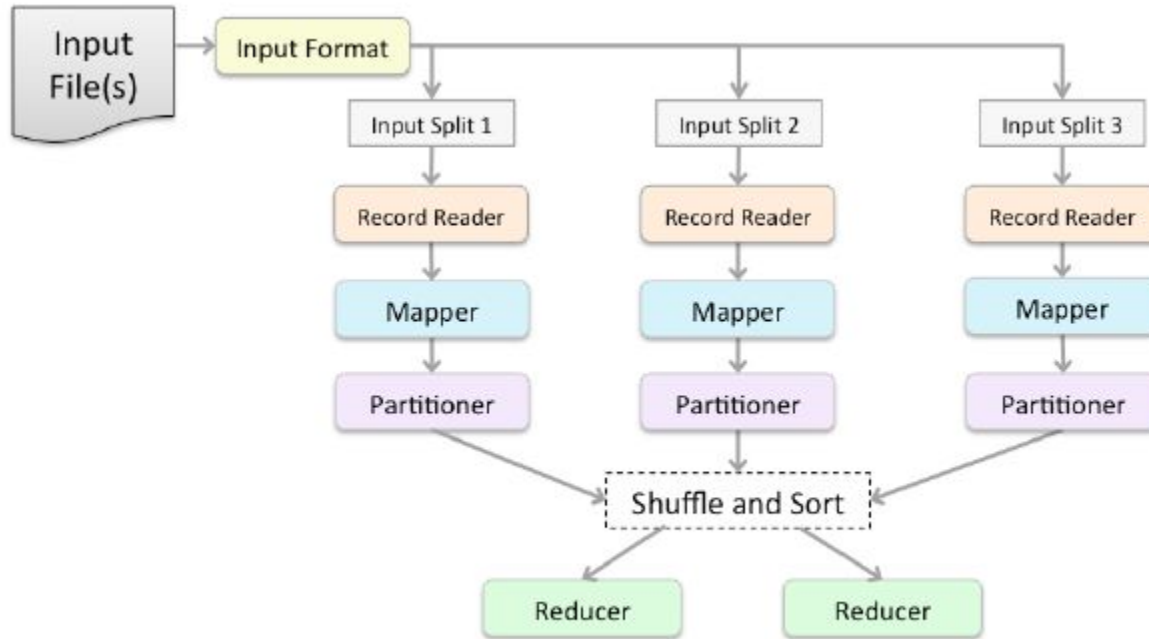
- Every **Split** is readable by a **RecordReader**
 - RecordReaders turn Splits into a iterable list of Records
- Every **MapTask** uses the **RecordReader** to feed Records to a **Mapper**
- Each **Mapper** processes one **Split**

MapReduce Processing



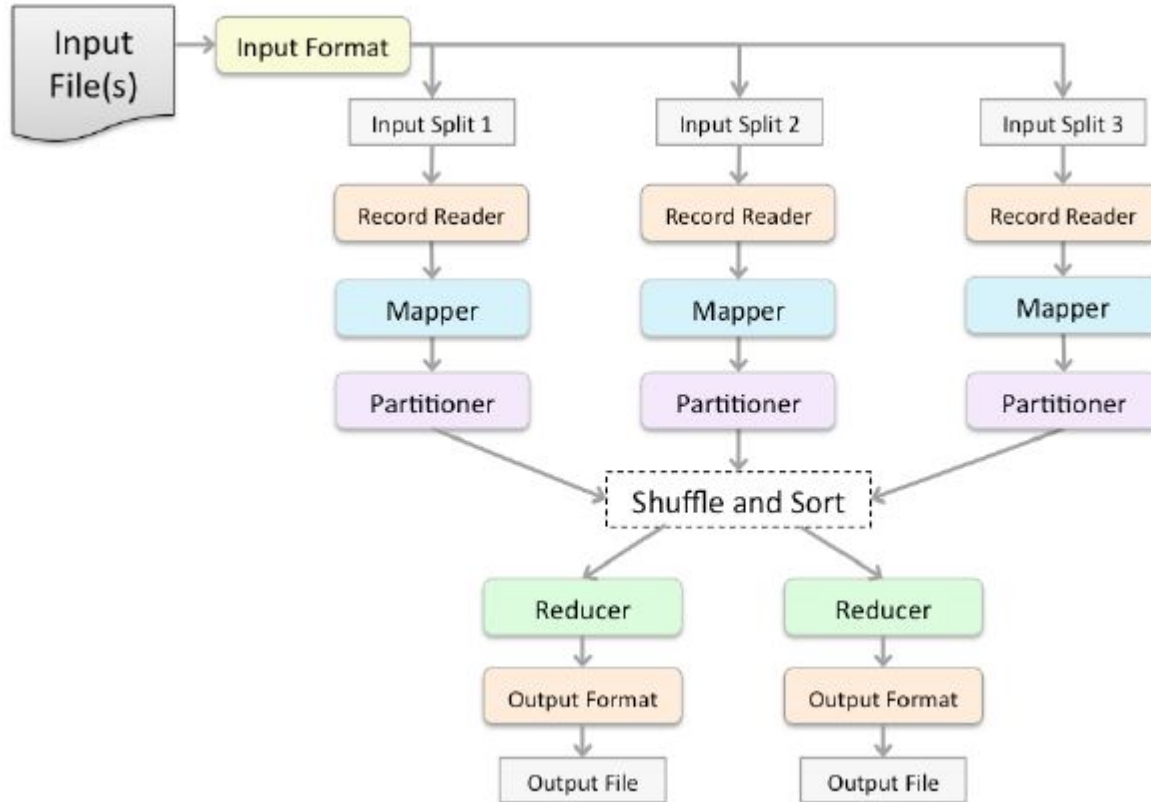
- **Mapper outputs are given a partition index**
- **Shuffle-sort (in the MapTask):**
 - **Mapper** results are sorted and **partitioned** into files
 - The **partitioned** files are available for HTTP fetch by Reducer Tasks

MapReduce Processing

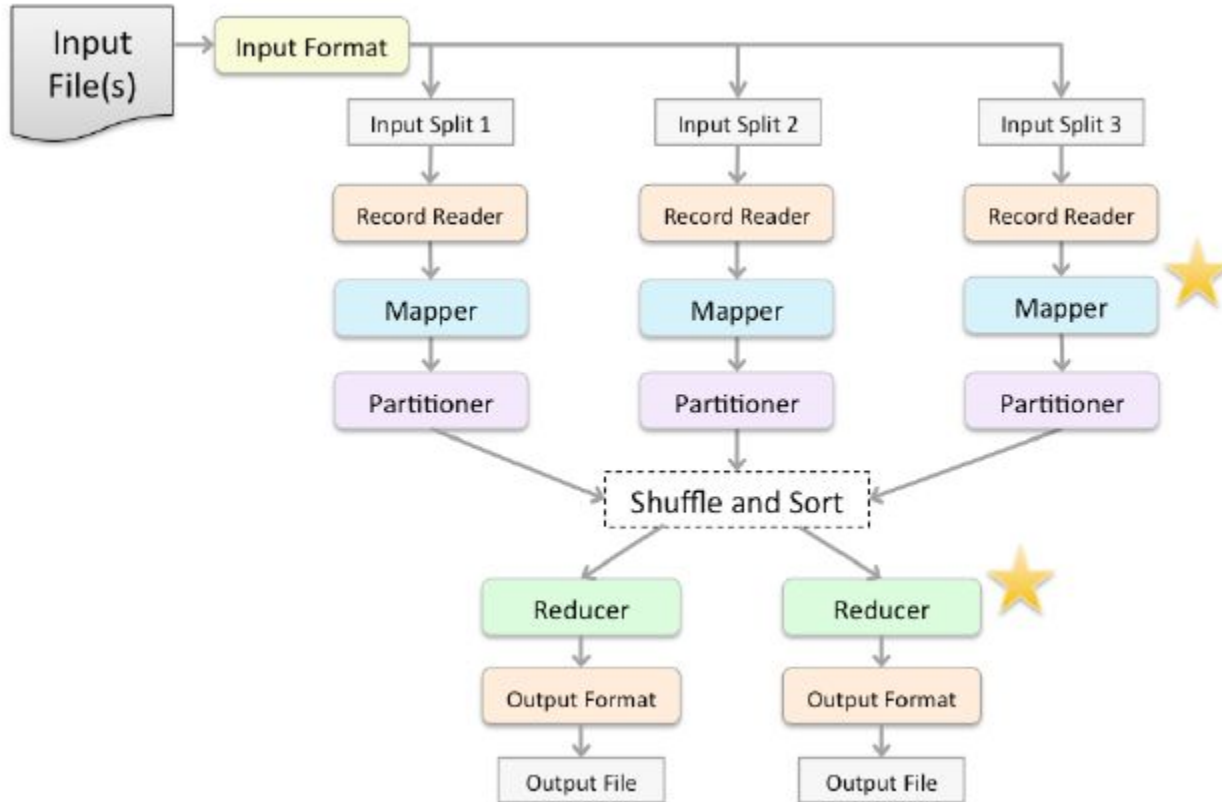


- **ShuffleSort** (In the ReduceTask) - The Reducertask fetches and merges partitioned files into a mergeFile
- The **Reducer** reads and processes the mergeFile

MapReduce Processing



MapReduce Processing




Review: MapTask steps

For each block of data, create a MapTask. Each MapTask does the following:

1. **InputFormat:** convert blocks into input splits that respect record boundaries.
2. **RecordReader:** parse split into records, write record as $\langle k, v \rangle$ pair.
3. **Mapper:** run function on each $\langle k, v \rangle$ pair and write $\langle k', v' \rangle$, as desired.
4. **Partitioner:** assigns an index, p , to each $\langle k, v \rangle$ pair $\Rightarrow \{p, k, v\}$

map-side
shuffle-sort

- 
5. **In the background, MapTask runs the spill-thread:**
 - aggregates all the values for a key into a list
 - partitions map output into buckets
 - writes each bucket to a different file, one for each reducer

Review: ReduceTask steps

reduce-side
shuffle-sort

1. **Network phase:** ReduceTasks fetch MapTask output files
 - Each ReduceTask is assigned a partition
 - A ReduceTask fetches the buckets for its partition from the MapTasks
 - Reduce tasks do not start 'reducing' until **all** the MapTasks are done.
2. **Merge-sort:** The ReduceTask merge-sorts all fetched buckets into one file
3. **Reducer:** processes each $\langle k, \{v_1, v_2, v_3 \dots\} \rangle$ and outputs results: $\langle \text{key}, \text{result} \rangle$
4. **OutputFormat:** determines format and writes the output. Usually, each Reducer produces a separate file.

Example: Average word length

(Covered in more detail in Practice 3)

- **Read in Shakespeare plays**
- **Mapper: find the first letter of each word**
 - write out the first letter as key
 - write out the length of the word as value
- **Reducer: read in the letter and list of lengths**
 - iterate through the list, summing
 - calculate the average
 - output the key (letter) and the average length

AvgWordLength - main method

```
public static void main(String[] args) throws Exception {
```

```
    Job job = Job.getInstance();  
    job.setJobName("Average Word Length");  
    job.setJarByClass(AvgWordLength.class);
```

Create and name the job
Find the jar containing this class

```
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

Define input and
output locations

```
    job.setMapperClass(LetterMapper.class);  
    job.setReducerClass(AverageReducer.class);
```

Set the Mapper and
Reducer classes

```
    job.setMapOutputKeyClass(Text.class);  
    job.setMapOutputValueClass(IntWritable.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(DoubleWritable.class);
```

Set the output classes for
the Mapper and Job

```
    System.exit(job.waitForCompletion(true) ? 0 : 1);
```

```
}
```

LetterMapper

```
package averageWordLength.solution;

public class LetterMapper extends Mapper<LongWritable, Text, Text, IntWritable>{

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException,
        InterruptedException {

        String line = value.toString();

        for (String word : line.split("\\W+")) {
            if (word.length() > 0) {
                String letter = word.substring(0, 1);
                context.write(new Text(letter), new IntWritable(word.length()));
            }
        }
    }
}
```


Letter Mapper in action

Input split

Four score and seven years ago our fathers brought forth on this continent a new nation, conceived in liberty and dedicated to the proposition that all men are created equal

Key, value inputs

Key	Value
0	Four score and seven years ago our
34	fathers brought forth on this continent
74	a new nation, conceived in liberty
109	and dedicated to the proposition that

Letter Mapper

map output

F	4	n	3
s	5	n	5
a	3	c	9
s	5	i	2
y	5	l	7
a	3	f	7
o	3	b	7
		f	5
		o	2
		t	4
		c	10
		a	1

Map-side shuffle-sort

n	3
n	5
c	9
i	2
l	7
f	7
b	7

mapper output

F	4
s	5
a	3
s	5
y	5
a	3
o	3

merge-sort

add partition index

key	values	partition index
a	1,3,3,3,3	0
b	7	0
c	9,7	0
d	9	0
e	5	0
f	4,7	0
i	7	1
l	9	1
m	3	1
n	3, 5	1
o	3	1
p	11	1

partition 0

a	1,3,3,3,3
b	7
c	9,7
d	9
e	5
f	4,7

part-m-00000

partition 1

i	7
l	9
m	3
n	3, 5
o	3
p	11

part-m-00001

AverageReducer

```
public class AverageReducer extends Reducer<Text, IntWritable, Text, DoubleWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        long sum = 0, count = 0;

        for (IntWritable value : values) {

            sum += value.get();
            count++;
        }
        if (count != 0) {

            double result = (double)sum / (double)count;
            context.write(key, new DoubleWritable(result));
        }
    }
}
```

Merge-sort and Reduce (Reduce Task)

part-0.3

a	7,2,1,3,3
b	7
c	9,7
d	9
e	5
f	4,7

part-0.1:
partition 0 from MapTask 1

a	1,3,3,3,3
b	7
c	9,7
d	9
e	5
f	4,7

part-0.2

	1,2,2,3
	6,6
	5,3
d	11
e	7,5
f	9, 7, 3

merge sort

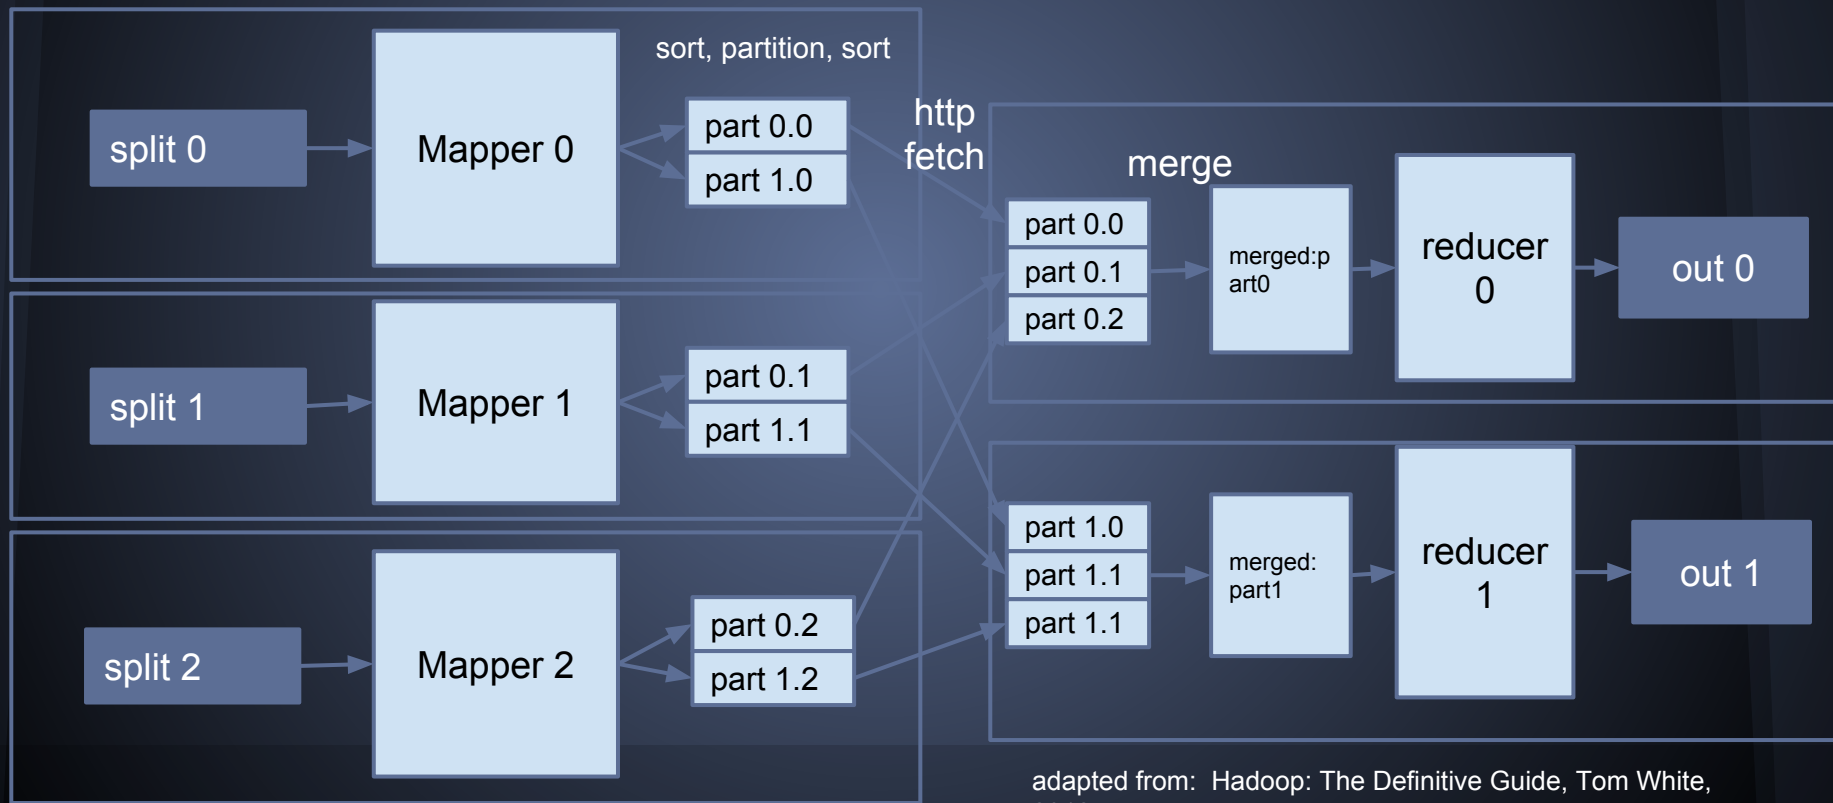
a	1,3,3,3,3,7,2,1,3,3,1,2,2,3
b	8,5,7,6,6
c	5,7,9,7,5,3
d	3,5,11
e	5,5,4,7,7,5
f	3,9,4,9,7,3

Average
Reducer
0

part-r-0000

a, 2.64
b, 6.4
c, 6.0
d, 6.33
e, 5.5
f, 5.83

OVERVIEW: map-reduce data flow



Creating and Running a MapReduce Job

Development

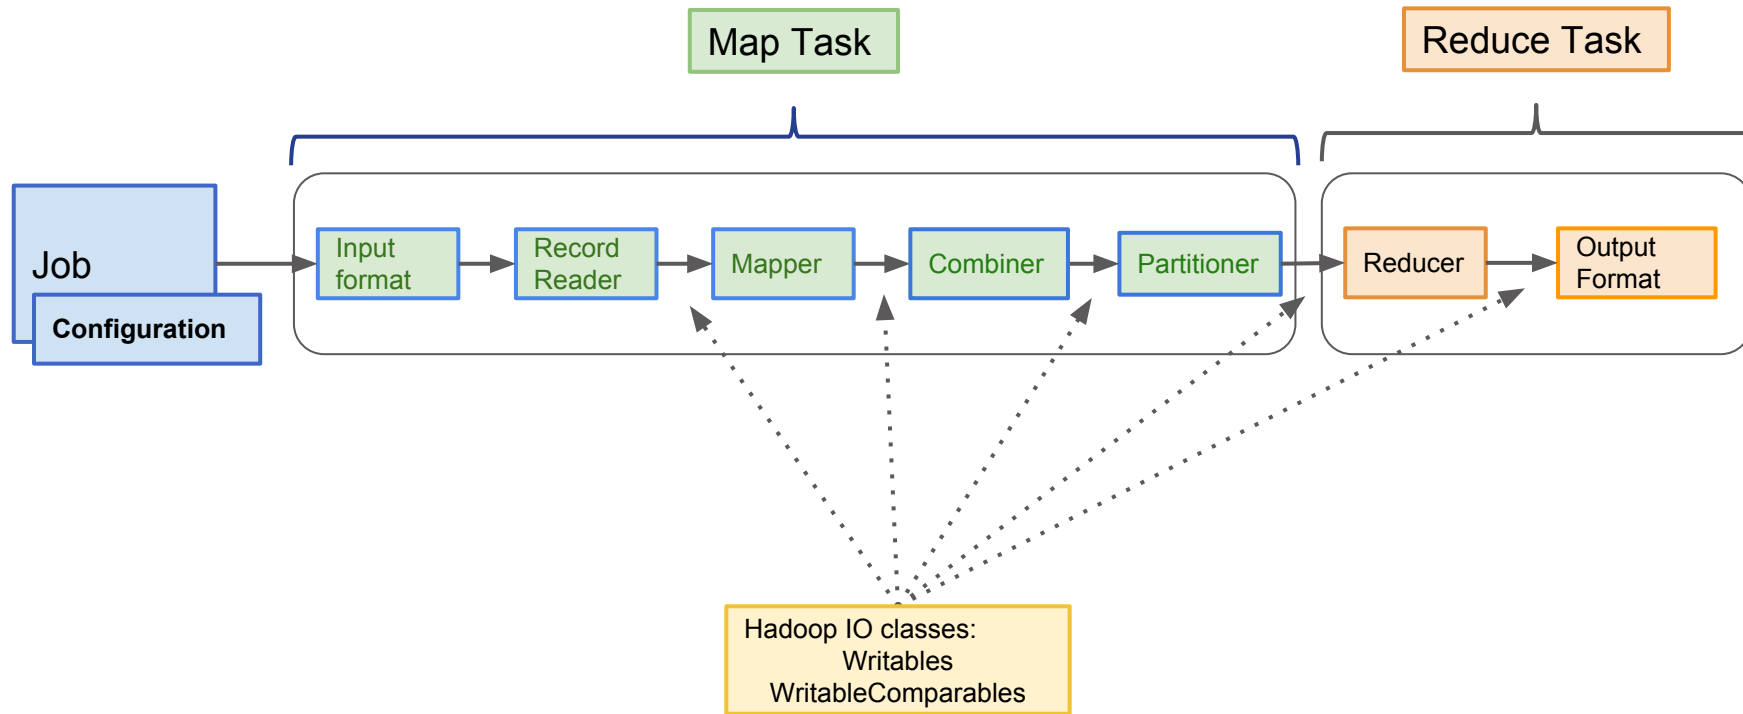
- Write the code in Eclipse
- Let Eclipse compile the code (automatic)
- Test the code by running it in Eclipse (use Run Configurations)

Submitting to the cluster


- Create a jar file from the compiled code
- Use the 'hadoop jar' command to submit

Detailed instructions in Practice 2

Important MapReduce Classes



Job configuration in the main method



The main method submits a job

- The main method creates a Job.
- The Job wraps a Configuration object.
- We define a Job using setters in the Job class.
 - set configuration properties
 - define the MR classes the developer has extended
 - define IO paths and IO types

**Going deeper: The Job implements MRJobConfig
(see `org.apache.hadoop.mapreduce.MRJobConfig`)**

The simplest job driver

```
public static void main(String[] args) throws Exception {  
  
    Job job = Job.getInstance();  
    job.setJobName("Identity Job");  
    job.setJarByClass(IdentityDriver.class);  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

Note: No Mapper or Reducer classes defined→ Top-level Mapper and Reducers are used:


- the top-level Mapper and Reducer do not change the data.
- result is the same as the input data except it is sorted (it went through shuffle-sort)

Job setters and their defaults

Setter	Default value
setInputFormatClass	TextInputFormat
setMapperClass	Mapper
setMapOutputKeyClass	LongWritable
setMapOutputValueClass	Text
setPartitionerClass	HashPartitioner
setNumReduceTasks	1
setReducerClass	Reducer
setOutputKeyClass	LongWritable
setOutputValueClass	Text
setOutputFormatClass	TextOutputFormat

If you do not use the setter, then the default is used.

Finding and formatting input data



Specifying input locations

- **Set the input location using InputFormat.**

```
FileInputFormat.setInputPaths(job, new Path(<dir>))
```

- **This will read the files in <dir> and execute them in MapReduce.**

- Won't read files that start with "." or "_" (hidden files)
- Can use wildcards to restrict input: /2010/*/Jan/*
- <dir> can be a directory or a file

- **To add multiple paths:**

```
FileInputFormat.addInputPath(job, new Path(<file>))
```


What is an InputFormat?

- Holds the location of the input, usually a file or directory.
- Formats data blocks into splits and then splits into records.
- InputFormats and OutputFormats are part of Hadoop IO
 - Hadoop IO is used by *both MR2 and Spark*
 - Details on InputFormats are covered in a later lecture

For now, we'll just be using the default **TextInputFormat**

Format type	Structured	Comments
Text files	No	Plain old text files. Records are assumed to be one per line.
Key-value text	Semi	Format for key-value pairs. Delimiters can be configured.
Whole file input		Processes one file per split.
Sequence File	Yes	A compressible format for binary data. Header holds metadata about contents. Has sync points for splitting.
Avro File	Yes	A flexible, compressible format with associate schema for handling complex, structured, evolving data.
DB (SQL statement)	Yes	A format for SQL statements from databases. (Most database tables are loaded using Sqoop)
Other: JSON, CSV, XML	Semi	Read in using TextFileFormat; parse in map using json/XML/csv specific libraries. May also use Hadoop streaming.

Mappers



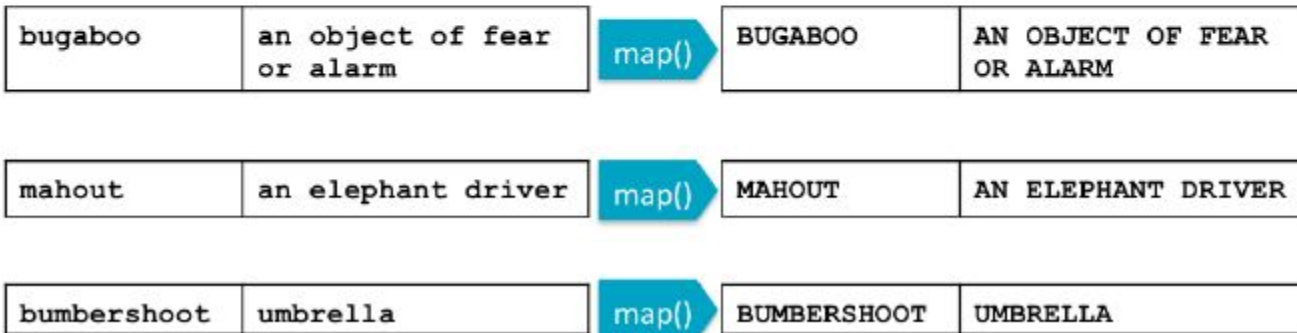
The Mapper runs on every MapTask

- **Hadoop attempts to ensure that Mappers run on nodes which hold their portion of the data locally, to avoid network traffic**
 - Multiple Mappers run in parallel, each processing a portion of the input data
- **The Mapper reads data in the form of key/value pairs**
 - The Mapper may use or completely ignore the input key
 - For example, a standard pattern is to read one line of a file at a time
 - The key is the byte offset into the file at which the line starts
 - The value is the contents of the line itself
 - Typically the key is considered irrelevant
- **If the Mapper writes anything out, the output must be in the form of key/value pairs**

Example Mapper: Upper Case Mapper

- Turn input into upper case (pseudo-code):

```
let map(k, v) =  
    emit(k.toUpper(), v.toUpper())
```



UpperCaseMapper

```
public class UpperCaseMapper extends Mapper<Text, Text, Text, Text> {  
  
    private Text uKey = new Text();  
    private Text uValue = new Text();  
  
    @Override  
    public void map(Text key, Text value, Context context) {  
  
        uKey.set(key.toString().toUpperCase());  
        uValue.set(value.toString().toUpperCase());  
        context.write(ukey, uvalue);  
  
    }  
  
}
```

Example Mapper: 'Explode' Mapper

- Output each input character separately (pseudo-code):

```
let map(k, v) =  
  foreach char c in v:  
    emit (k, c)
```

pi	3.14
----	------

map()

pi	3
pi	.
pi	1
pi	4

145	kale
-----	------

map()

145	k
145	a
145	l
145	e


ExplodeMapper

```
public class ExplodeMapper extends Mapper<Text, Text, Text, Text> {  
  
    private Text c = new Text();  
  
    @Override  
    public void map(Text key, Text value, Context context) {  
  
        char[] array = value.toString().toCharArray();  
        for (int i = 0; i < array.length; i++)  
            context.write(key, c.set(array[i]));  
  
    }  
  
}
```

Example Mapper: 'Filter' mapper

- Only output key/value pairs where the input value is a prime number (pseudo-code):

```
let map(k, v) =  
  if (isPrime(v)) then emit(k, v)
```



FilterMapper

```
public class FilterMapper extends Mapper<Text, IntWritable, Text, IntWritable> {  
  
    @Override  
    public void map(Text key, IntWritable value, Context context) {  
  
        if (isPrime(value))  
            context.write(key, value);  
  
    }  
  
    private boolean isPrime(IntWritable value) {  
        ... your sieves here ...  
    }  
  
}
```

Example Mapper: Changing Keyspaces

- The key output by the Mapper does not need to be identical to the input key
- Example: output the word length as the key (pseudo-code):

```
let map(k, v) =  
    emit(v.length(), v)
```



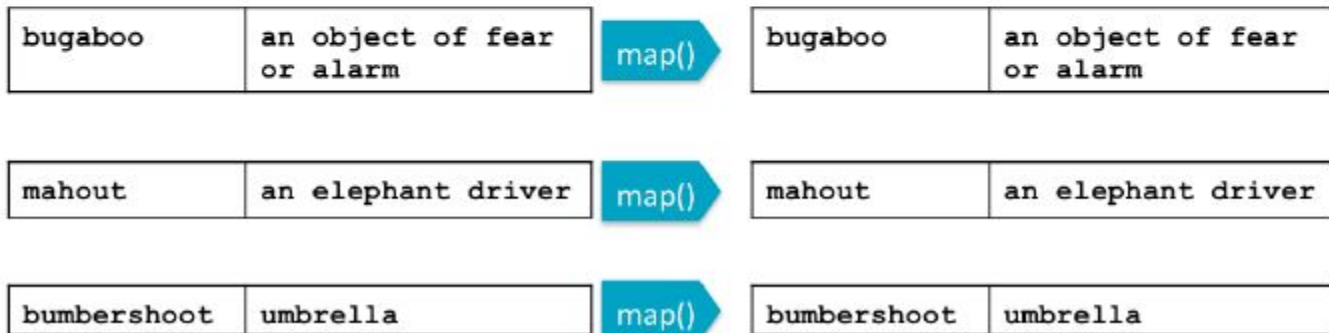
Changing Key spaces: LengthMapper

```
public class LengthMapper extends Mapper<IntWritable, Text, IntWritable, Text> {  
  
    IntWritable lengthKey = new IntWritable();  
  
    @Override  
    public void map(IntWritable key, Text value, Context context) {  
  
        lengthKey.set(value.toString().length());  
        context.write(lengthKey, value);  
  
    }  
  
}
```


Example Mapper: Identity Mapper

- Emit the key,value pair (pseudo-code):

```
let map(k, v) =  
    emit(k, v)
```



Reducers



For each input key, the Reducer *reduces* the list of values to a smaller set of values.

The Reducer

- **After the Map phase is over, all intermediate values for a given intermediate key are combined together into a list**
- **This list is given to a Reducer**
 - There may be a single Reducer, or multiple Reducers
 - All values associated with a particular intermediate key are guaranteed to go to the same Reducer
 - The intermediate keys, and their value lists, are passed to the Reducer in sorted key order
- **The Reducer outputs zero or more final key/value pairs**
 - These are written to HDFS
 - In practice, the Reducer usually emits a single key/value pair for each input key

Example Reducer: Sum Reducer

- Add up all the values associated with each intermediate key (pseudo-code):

```
let reduce(k, vals) =  
  sum = 0  
  foreach int i in vals:  
    sum += i  
  emit(k, sum)
```

the	1
	1
	1
	1



the	4
-----	---

SKU0021	34
	8
	19



SKU0021	61
---------	----

SumReducer code

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```


Example Reducer: Average Reducer

- Find the mean of all the values associated with each intermediate key (pseudo-code):

```
let reduce(k, vals) =  
  sum = 0; counter = 0;  
  foreach int i in vals:  
    sum += i; counter += 1;  
  emit(k, sum/counter)
```

the	1
	1
	1
	1



the	1
-----	---

SKU0021	34
	8
	19



SKU0021	20.33
---------	-------

Average Reducer code

```
public class AvgReducer extends Reducer<IntWritable, IntWritable, IntWritable, DoubleWritable> {  
  
    DoubleWritable average = new DoubleWritable(); ← create a holder for average  
  
    @Override  
    protected void reduce(IntWritable key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
  
        int sum = 0;  
        int count = 0;  
        for (IntWritable value : values) {  
            sum += value.get();  
            count++;  
        }  
        average.set(sum / (double) count);  
        context.write(key, average);  
    }  
}
```

Iterate through the values in the list, adding to the sum and incrementing the counter.

value.get() retrieves the integer.
average.set() sets the double.

Example Reducer: Identity Reducer

- The Identity Reducer is very common (pseudo-code):

```
let reduce(k, vals) =  
  foreach v in vals:  
    emit(k, v)
```

bow	a knot with two loops and two loose ends
	a weapon for shooting arrows
	a bending of the head or body in respect



bow	a knot with two loops and two loose ends
bow	a weapon for shooting arrows
bow	a bending of the head or body in respect

28	2
	2
	7



28	2
28	2
28	7

Caveat: keep track of types

the output types of the mapper
must match
the input types of the reducer

Mapper and Reducer outputs must match main settings

```
public static void main(String[] args) throws Exception {
```

```
    job.setMapOutputKeyClass(Text.class);  
    job.setMapOutputValueClass(IntWritable.class);
```

```
    ...
```

```
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(FloatWritable.class);
```

job defined in main

```
public class MaxTempMapper extends Mapper<LongWritable, Text, Text, IntWritable>
```

```
public class MaxTempReducer extends Reducer<Text, IntWritable, Text, FloatWritable>
```


Mapper output must match Reducer input

```
public class MaxTempMapper extends Mapper<LongWritable, Text, Text, IntWritable>
```

Map outputs must match Reducer inputs.

```
public class MaxTempReducer extends Reducer<Text, IntWritable, Text, IntWritable>
```

Writing formatted output data



Specifying output locations

- **define the output location using `OutputFormat`:**

```
FileOutputFormat.setOutputPath(job, new Path(<dir>))
```

- **This defines the directory that receive the final (reduced) results.**
- **This directory must not exist - MapReduce will create it.**

Output formats

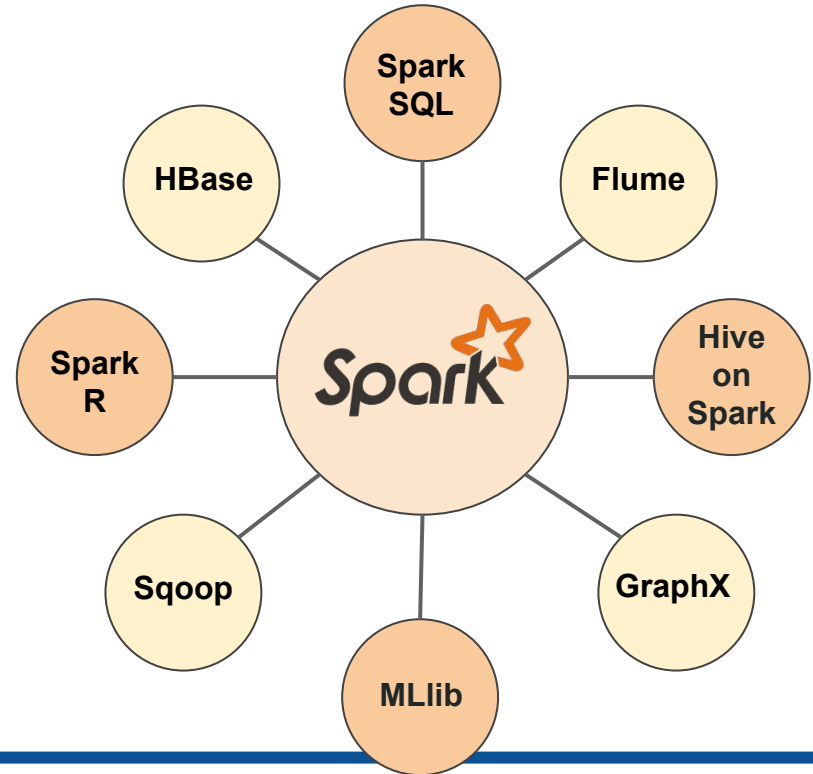
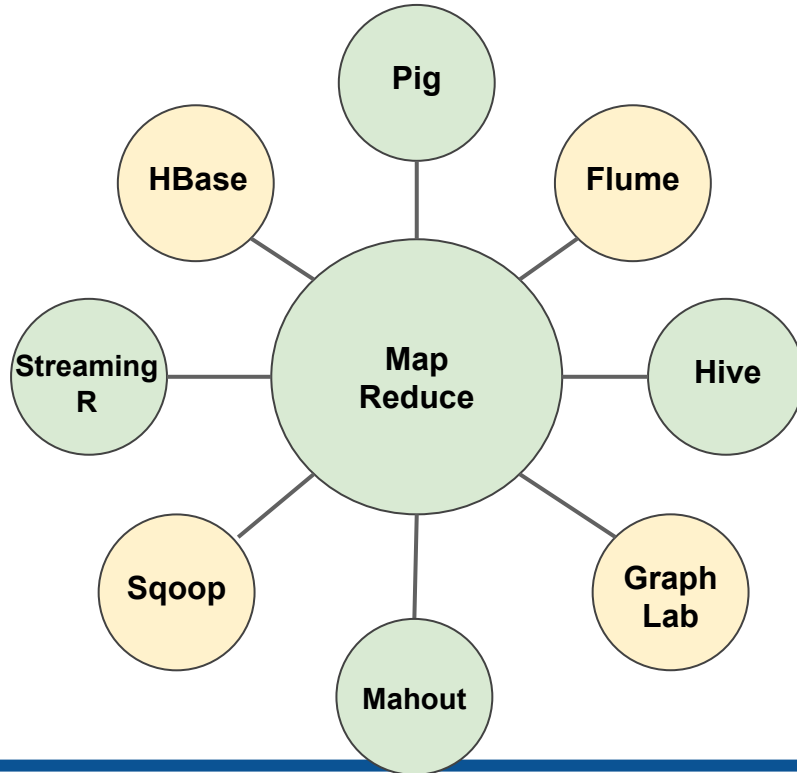
- **For now, we will use the default for the `OutputFormat`.**

```
TextOutputFormat.class;
```

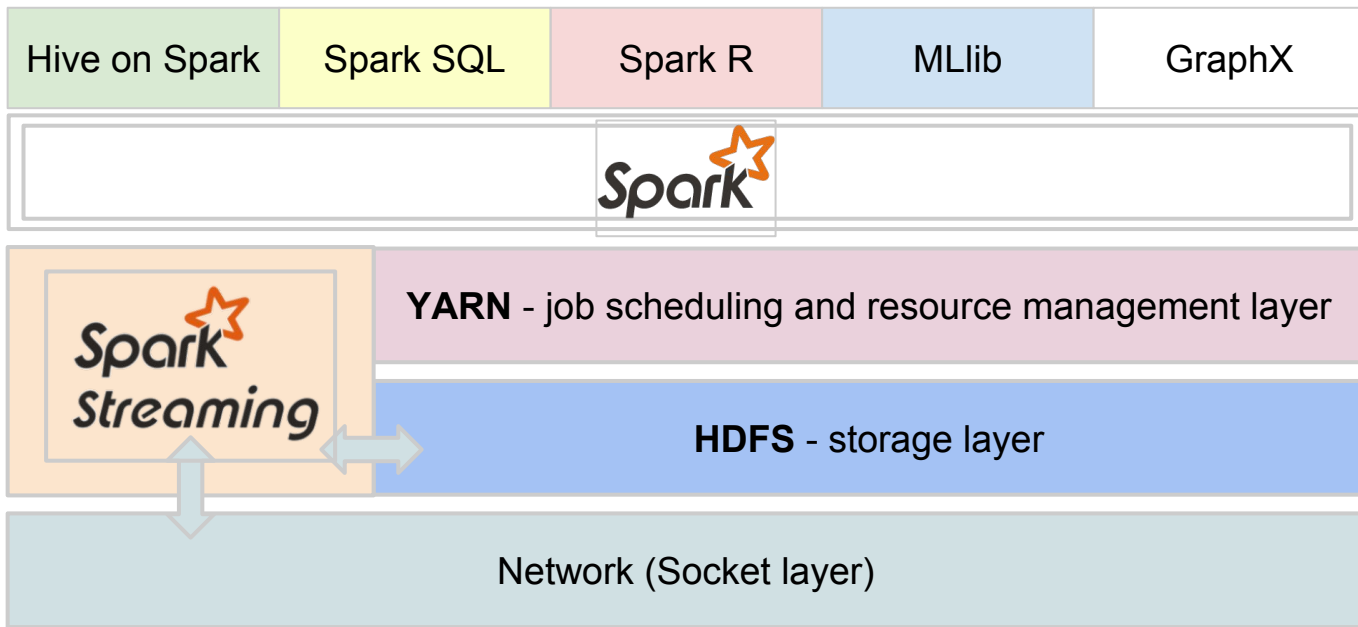
- Regardless of the output classes used, it will write out the results as Text (String)
- Very forgiving...

Spark

<https://cwiki.apache.org/confluence/display/SPARK/Committers>



Spark Analytics





- Founded in late 2013
- by the creators of Apache Spark (Matei Zaharia's PhD dissertation)
- Original team from UC Berkeley AMPLab
- Raised \$47 million in 2 rounds
- <100 employees, 100% recommend on Glassdoor
- They're hiring! (<https://databricks.com/company/careers>)
- Contributed more than 75% of the code in Spark

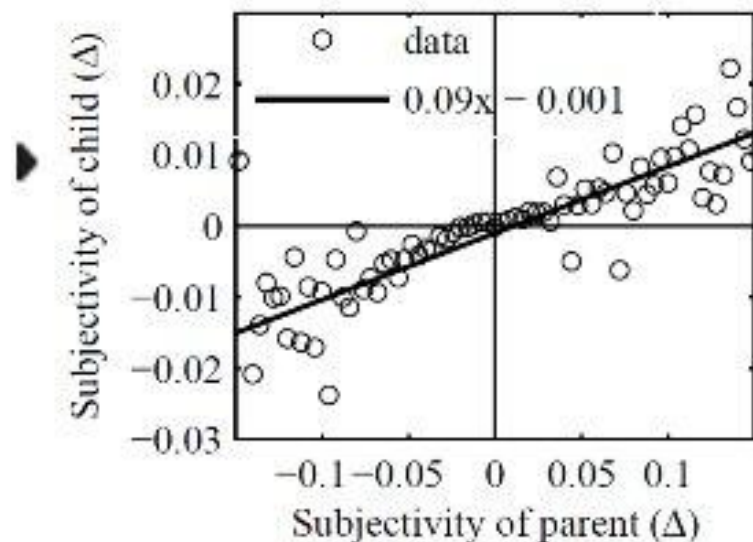
There's hype...

 **hadoop** + Spark  =


**A Winning
Combination**

The course of sentiment

- Cascades “heats” up early and then cool off



Subjectivity of the child and the parent are correlated. Sentiment flows!



There's change: Spark growth

Core:


- **2010: 14,000 LOC**

- **2012: 20,000 LOC**

- **2014: 50,000 LOC**

- **2016: 70,000+ LOC**

- including new libraries: 300,000+ LOC



What is Spark?

- A distributed in-memory compute system
- Can use Hadoop/YARN
- Uses the Map Reduce paradigm

Load -> Split -> Map -> Partition -> Shuffle-sort -> Reduce -> Output

- Read/write to HDFS
- Uses Hadoop IO (input and output formats, writables)

What is Spark?

- High-level functionality (joins, aggregates, group by, filter)
- Amazing job choreography
 - MR2 can only execute two tasks, in order: Map and Reduce
 - Spark can create a jobs executing many data transformations
 - While a complex problem might require several MR2 jobs, Spark can execute the same problem in one job.
- Spark's succinct code can represents complex jobs
 - No need for Cascading or Oozie

MR2 tasks vs Spark transformations

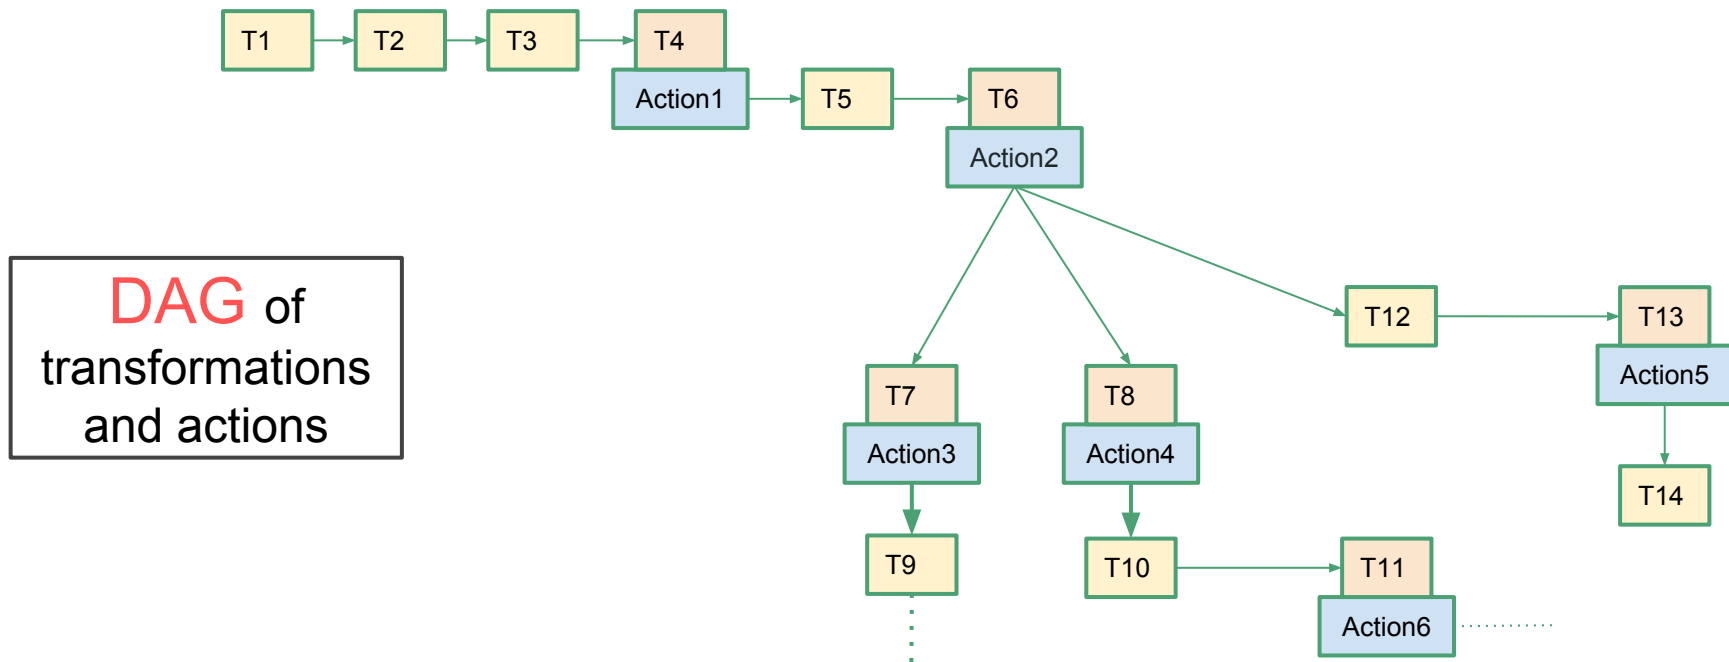
MR2 executes each job with just two tasks: MapTask and ReduceTask

Spark executes a job with many transformations.

- Narrow transformations are like MapTasks - they do not involve shuffle-sort
- Wide transformations are like ReduceTasks - they trigger a shuffle before the transformation begins.

Use case - Twitter

The processing pipeline for a Spark application with transformations (T) and Actions.



More about Spark's allure

- Solves the hard problems
 - Easily moves data from one MR job to another
 - Shuffles in-memory
 - Caches variables and recycles JVMs for tasks
- Solves user problems
 - Multiple languages
 - Python and Scala shell
 - Applications in Java, Python and Scala
 - Well-integrated with Spark R, Spark SQL, MLlib and GraphX
 - Easy data loads, powerful keywords, **concise**

Example: Word count (written in Scala)

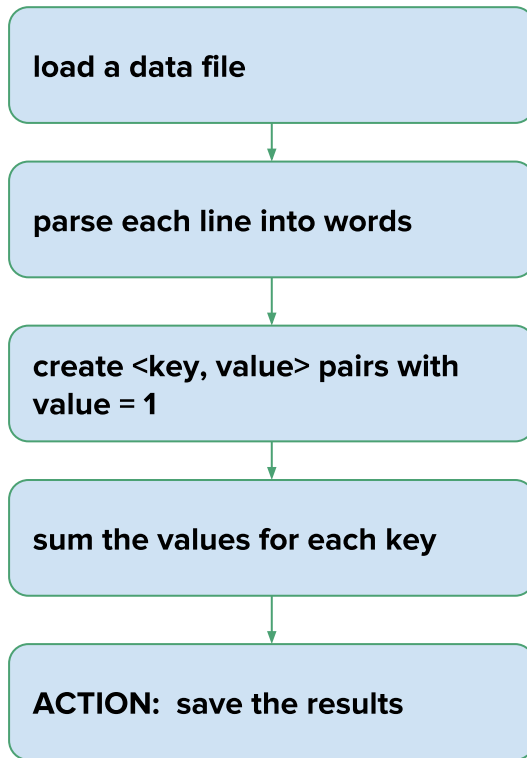
```
val text = sparkContext.textFile("file:///...")
```

```
val words = text.flatMap(line => line.split(" "))
```

```
val wordPairs = words.map(word => (word, 1))
```

```
val counts = wordPairs.reduceByKey(_ + _)
```

```
counts.saveAsTextFile("hdfs:/...")
```



Counting the number of times “love” appears in Shakespeare’s poetry

```
> val lines = sc.textFile("file:///home/cloudera/datasets/shakespeare/poetry")
```

```
> val words = lines.flatMap(line => line.split("\\W+"))
```

```
> val loveWords = words.filter(word=> word.contains("love"))
```

```
> val wordCount = loveWords.map(lword=> (lword, 1))
```

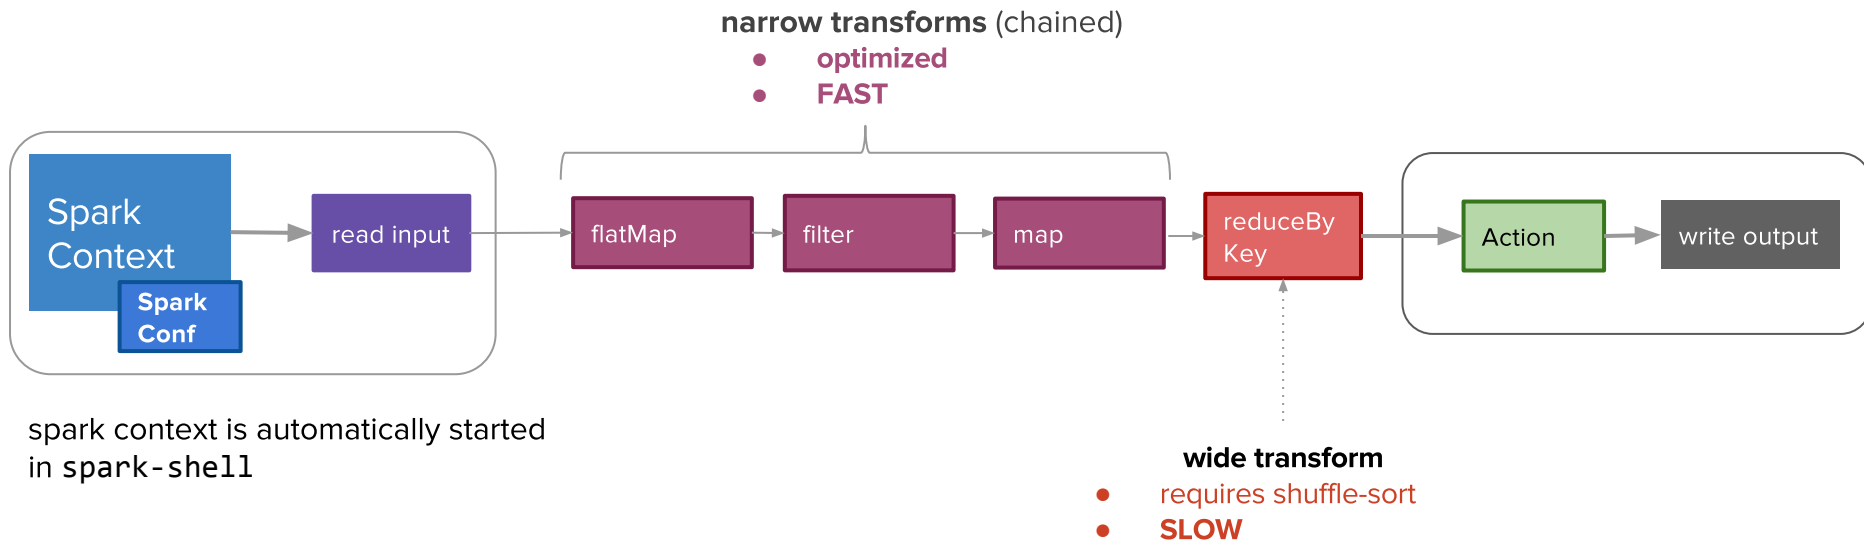
```
> val counts = wordCount.reduceByKey(_+_)
```

```
> counts.collect()
```

we could write
this in one
Mapper in MR2

this requires a
shuffle before we
start

Spark workflow for “love” count



At home practice

1. Start your VM
2. Open a terminal
3. Start the spark REPL:

```
$ spark-shell
```


4. Run the example by typing this

```
scala> val lines = sc.textFile("file:///home/cloudera/datasets/shakespeare/poetry")
```

```
scala> val words = lines.flatMap(line => line.split("\\W+"))
```

```
scala> val loveWords = words.filter(word => word.contains("love"))
```

```
scala> val wordCount = loveWords.map(word => (word, 1))
```

```
scala> val counts = wordCount.reduceByKey(_+_)
```

```
scala> counts.collect()
```

Stopping

5. stop the context

```
scala> sc.stop()
```

6. stop the REPL

```
scala> :quit
```

RDD - the core class of Spark

- compile-time type-safe
- lazy...
 - transform operations describe how to change the data
 - **map, filter, join, reduceByKey**
 - action operations actually start the processing and produce output
 - **take, count, first, foreach, collect**
- based on the Scala collections API
- most Spark RDD operations == Hive and MapReduce functionality

RDDs in the example

```
> val lines = sc.textFile("file:///home/cloudera/datasets/shakespeare/poetry")
```

lines is a file-backed RDD

```
> val words = lines.flatMap(line => line.split("\\W+"))
```

```
> val loveWords = words.filter(word => word.contains("love"))
```

```
> val wordCount = loveWords.map(word => (word, 1))
```

```
> val counts = wordCount.reduceByKey(_+_)
```

words, **loveWords**, **wordCount** and **counts** are RDDs derived from lines


```
> counts.collect()
```

collect() is an action - it starts processing and prints out **counts**


RDD operations

- **Two types of RDD operations**

- Actions – return values



- Transformations – define a new RDD based on the current one(s)



- **Pop quiz:**

- Which type of operation is `count()`?

Learning Spark

- Learn RDDs - focus on depth, learn the “how” of the processing
- Learn DataFrames and Datasets - focus on “use”, best for analysis
- Focus on Spark and Spark Streaming - for dataops
- Focus on Spark ecosystem (MLlib, SparkR, SparkSQL) - for analysis

Job execution

How the “count love words” example is processed

Job scheduling and execution

Very broad overview

Much more detail here - current, free book

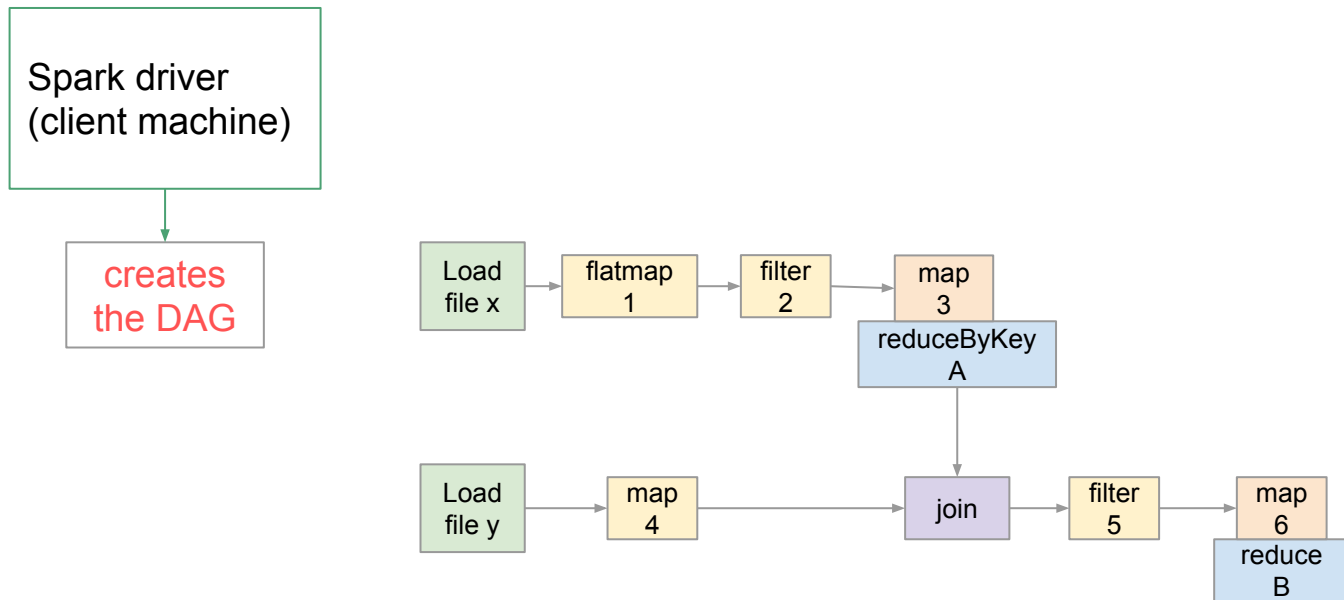
<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-dagscheduler.html>

Spark has a planning and execution engine

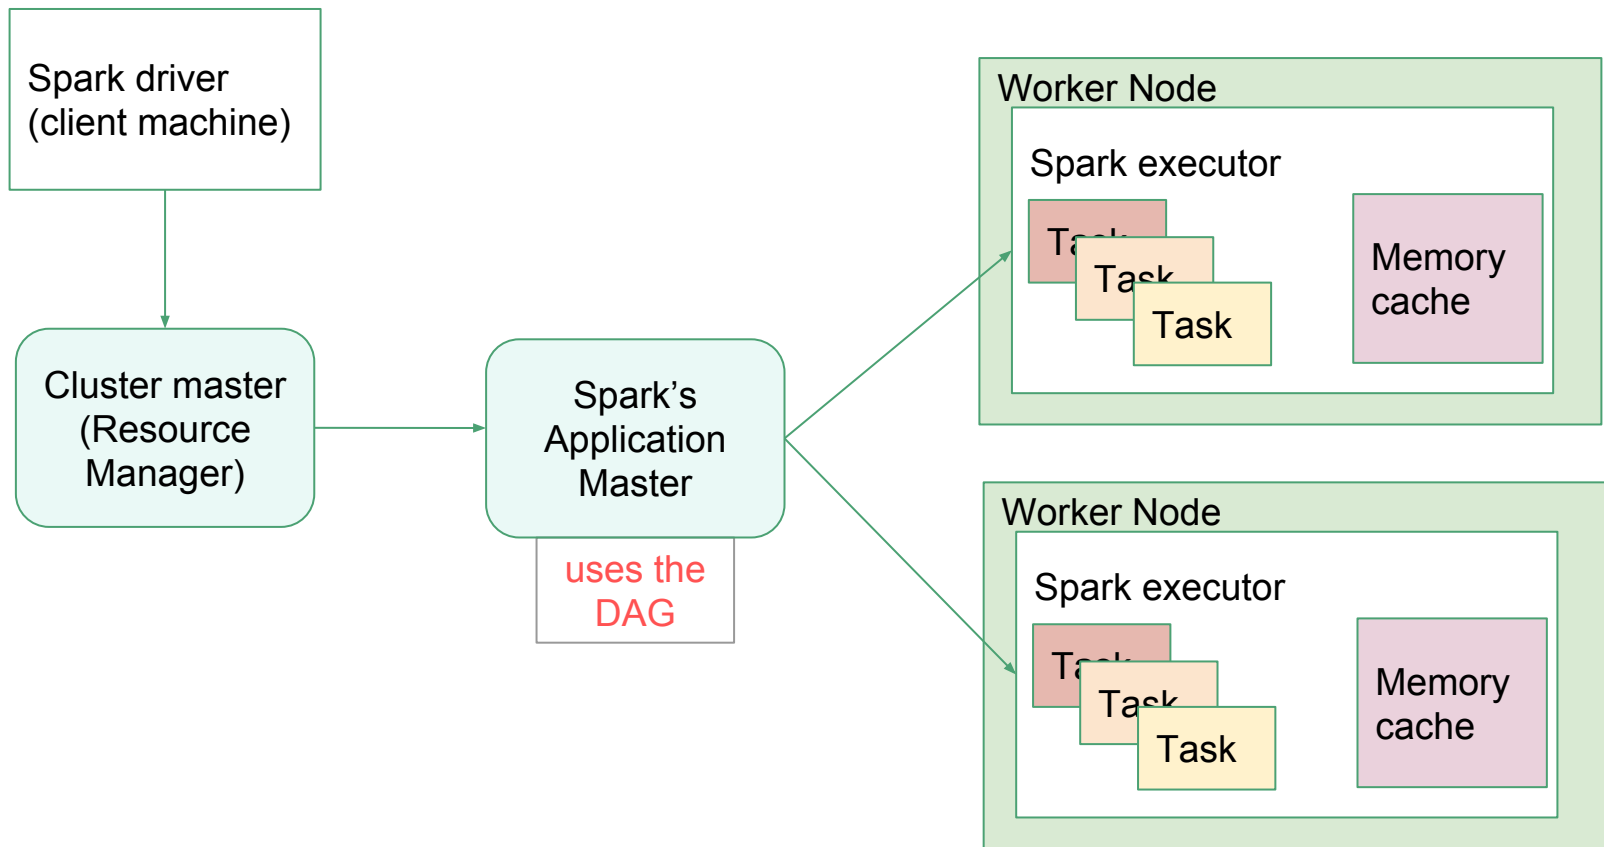
- Creates a plan for the job
- The job is executed in stages
- Narrow transformations are optimized (chained)
 - run on the same executor
 - use the same data
 - process the data line-by-line through the whole chain

Execution is managed by the DAGScheduler

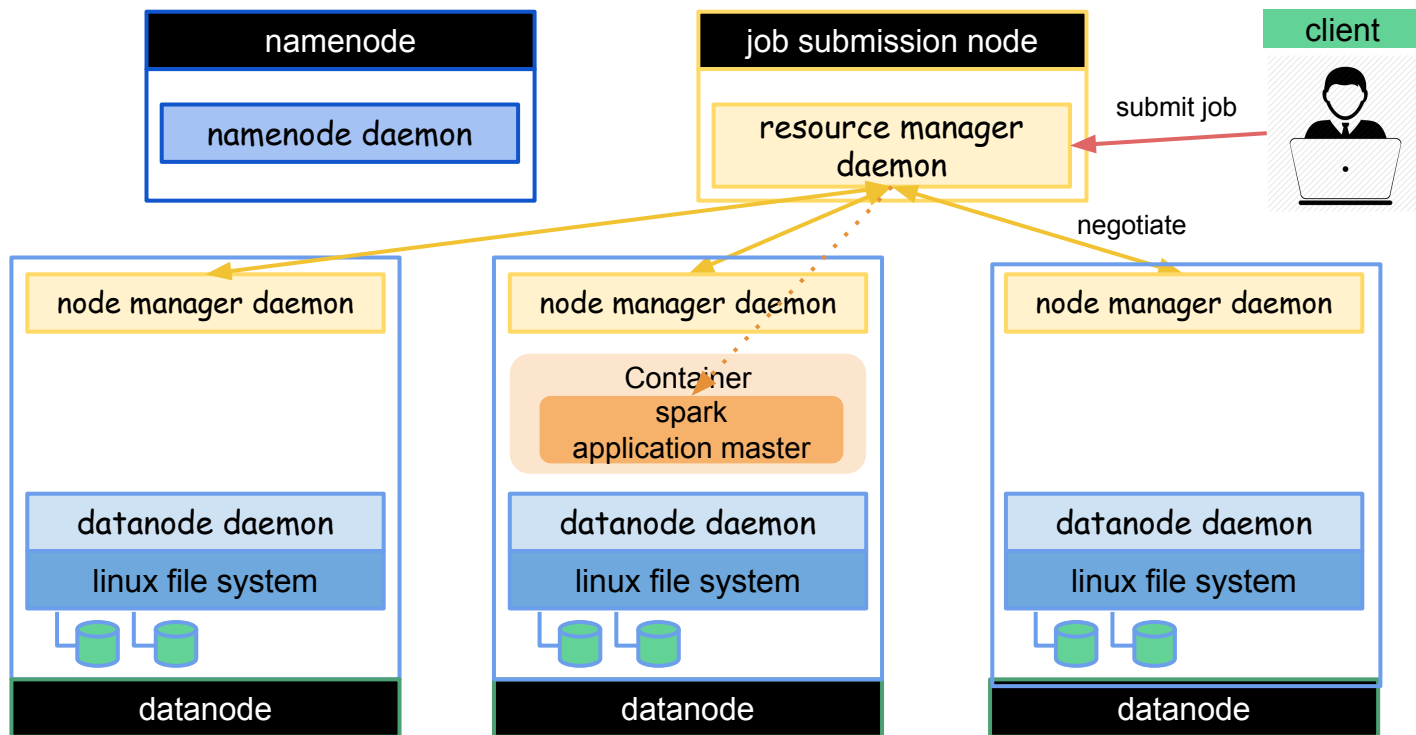
Spark execution(1)



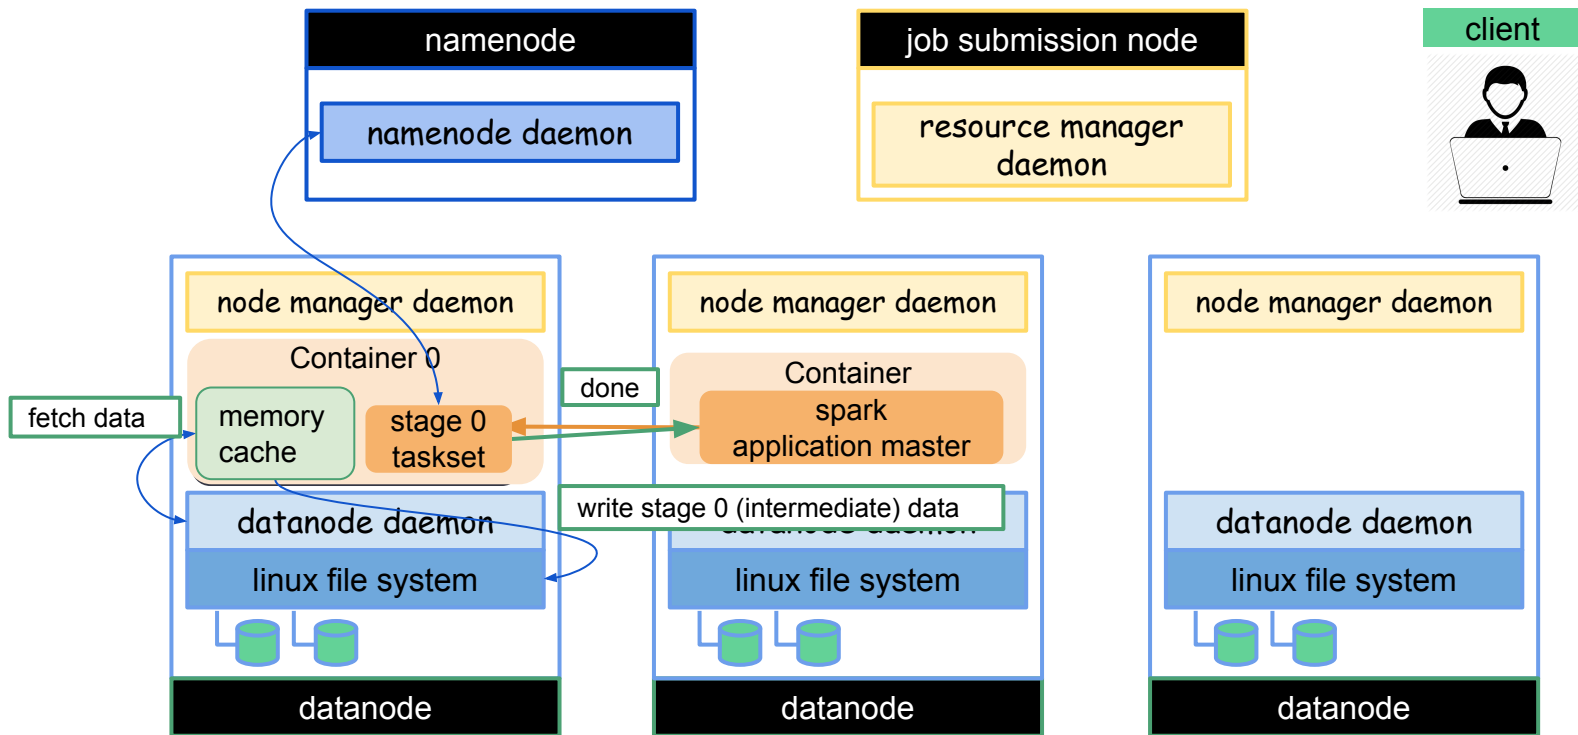
Spark execution(2)



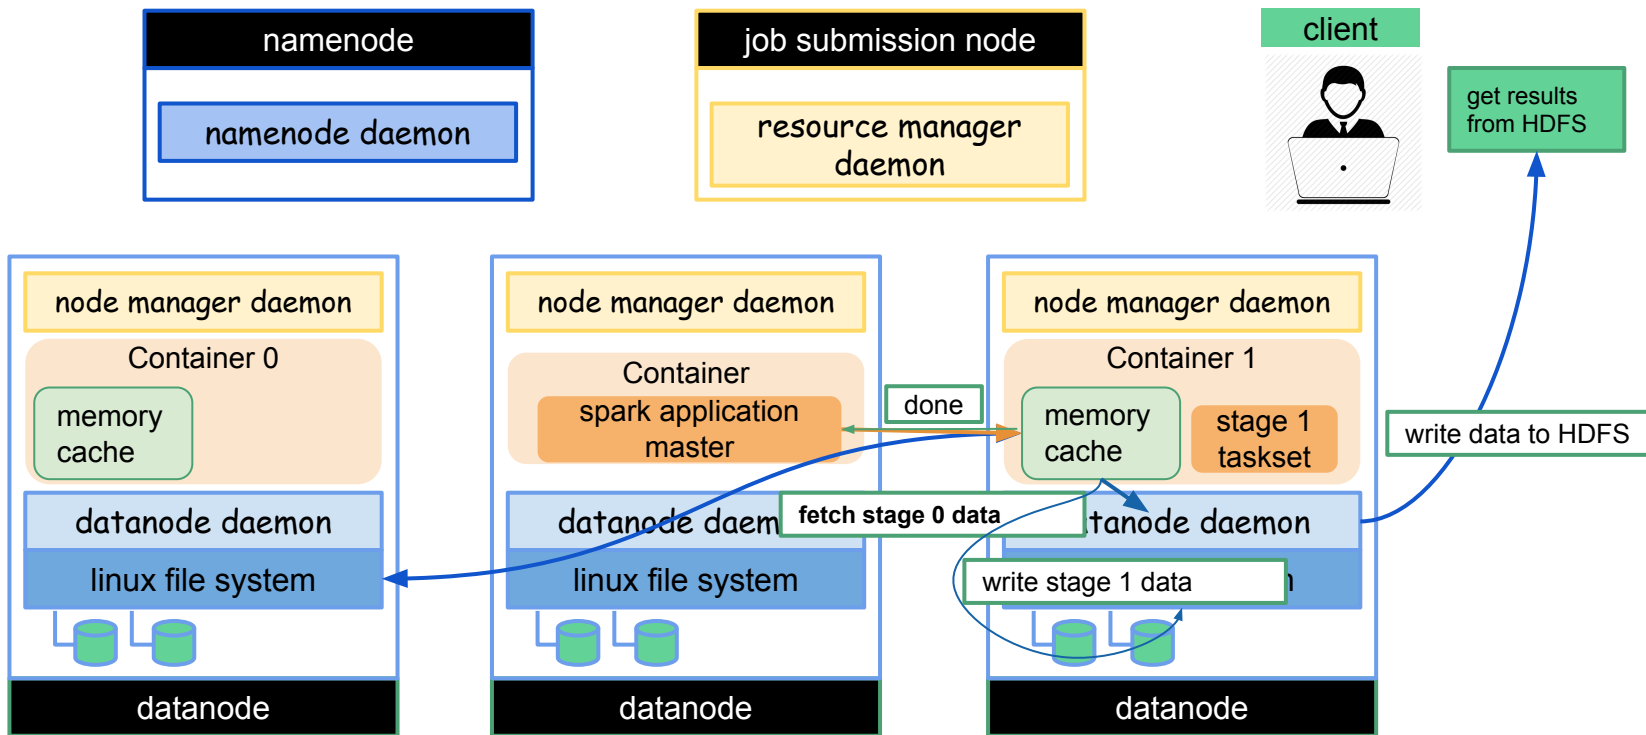
Running a 3-stage Spark job on YARN



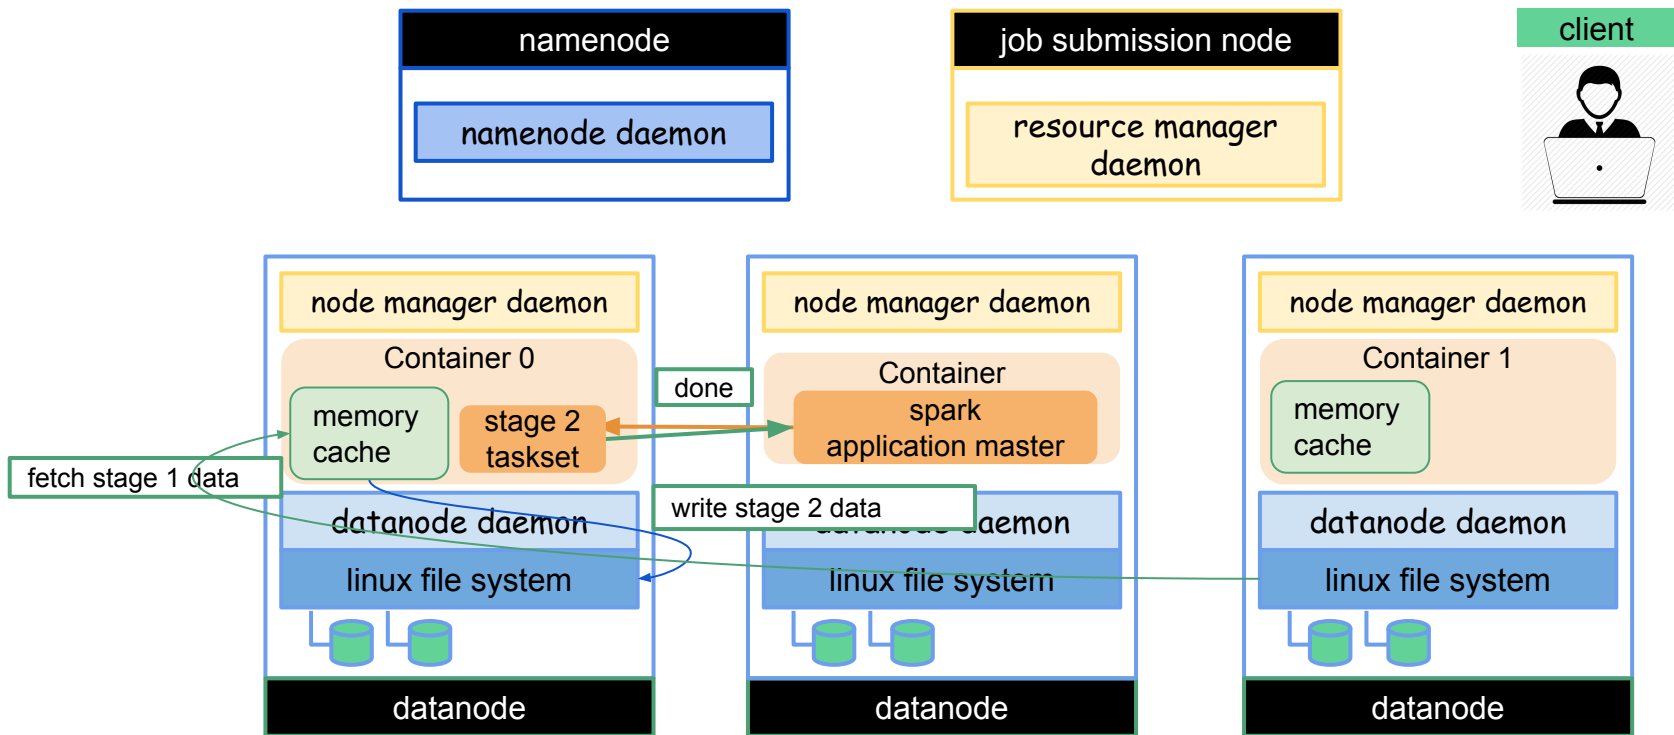
Running a 3-stage Spark job on YARN



Running a 3-stage Spark job on YARN

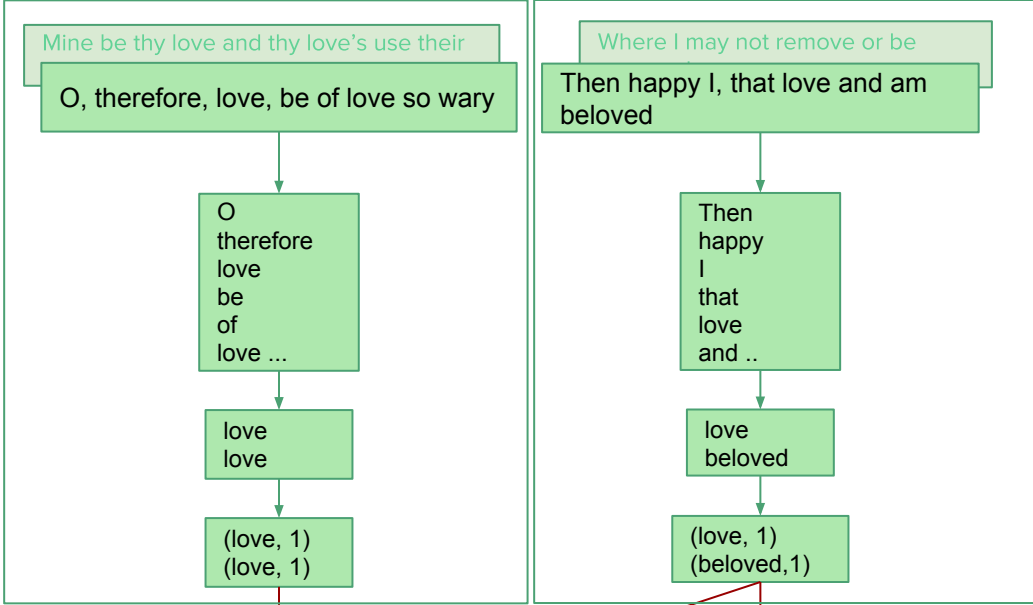
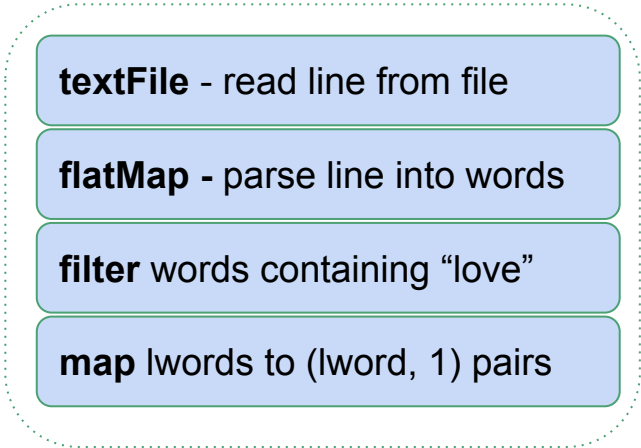


Running a 3-stage Spark job on YARN

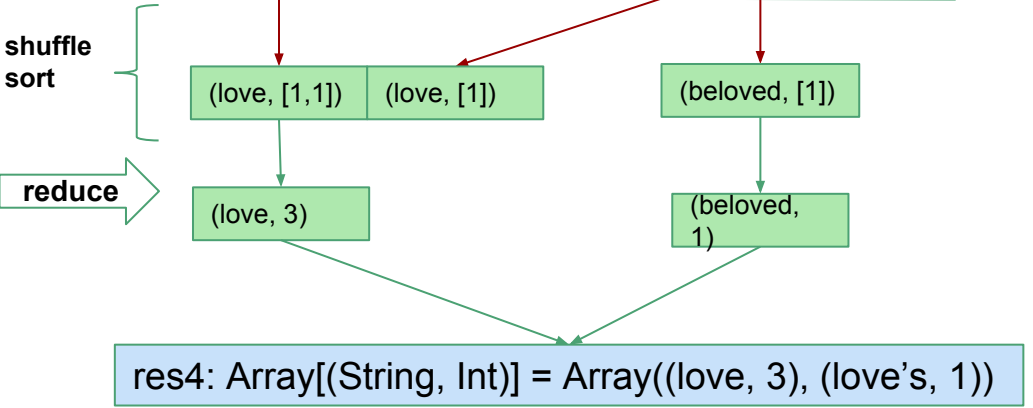
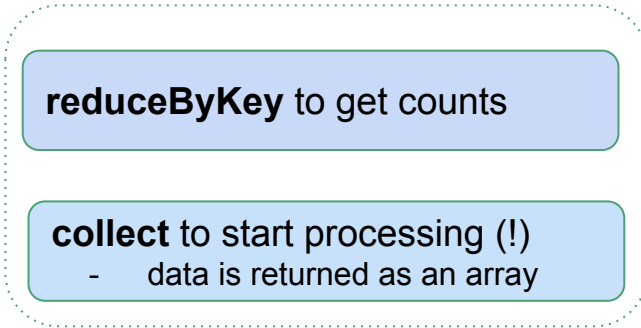


So, Spark jobs run in stages?

Stage 1

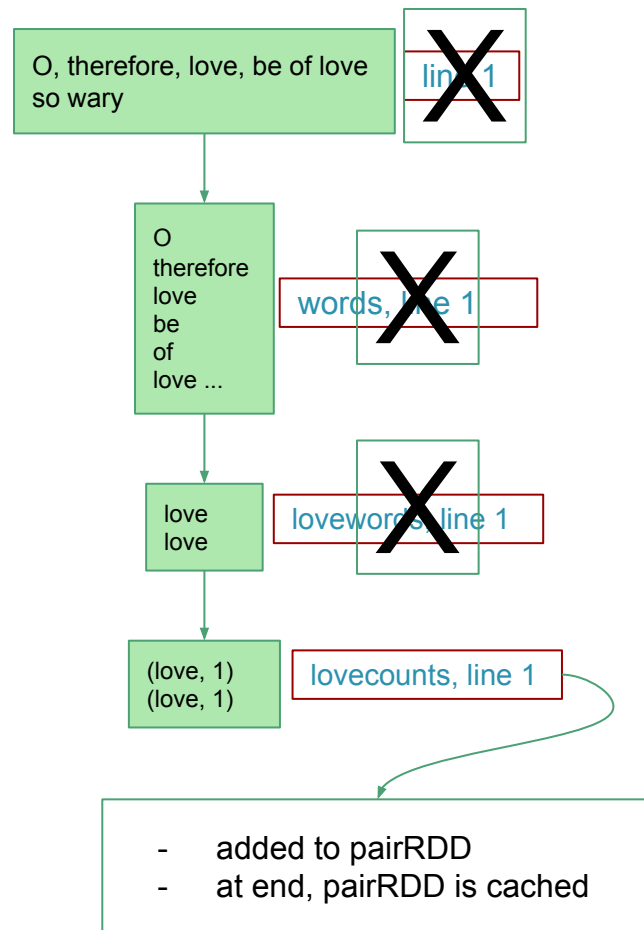


Stage 2



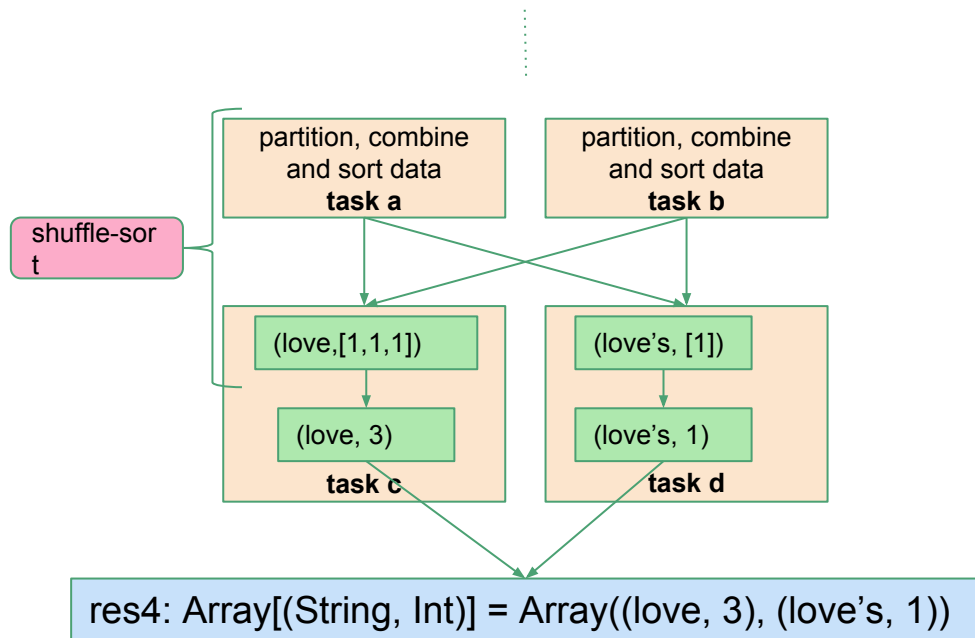
Stage 1: a chain of narrow transforms: *textFile -> flatmap -> filter -> map*

- Assign an executor to each input split
 - executes the chain of transforms
 - each line in a split is processed thru the chain
 - optimizes to chain to look/act like a single task
- Intermediate results (RDDs) are NOT stored
 - Records process consecutively
 - Each record is fed through all transforms
 - This is called “pipelining”
- At the end of a stage, the data is cached.
 - If we want to “save” an intermediate RDDs, use `<name of rdd>.cache()`



Stage 2:

- a wide transformation:
reduceByKey
 - performs shuffle-sort
 - runs the *ReduceByKey* function
- at end, **collect** prints results:
 - outputs an iterator (e.g. for an array)
 - used for small results




Summary

```
val lines = sc.textFile("/home/cloudera/datasets/shakespeare")  
val words = lines.flatMap(line => line.split("\\W+"))  
val loveWords = words.filter(word=> word.contains("love"))  
val loveCounts = loveWords.map(lword=> (lword, 1))
```

Stage 1

```
val counts = loveCounts.reduceByKey(_+_)  
counts.collect()
```

Stage 2



That collect() action - huh?

- collect() is an **action**
 - *starts* the processing by contacting DAGScheduler
 - acts like: `job.waitForCompletion()` in MapReduce
 - DAGScheduler is an event planner
 - everything is planned and ready to go
 - just needs to be told when to start the party
 - Example: “`counts.collect()`” was the **action** that started the job
 - Remember, you saw processing after “`counts.collect()`”

Important differences: traditional MR and Spark

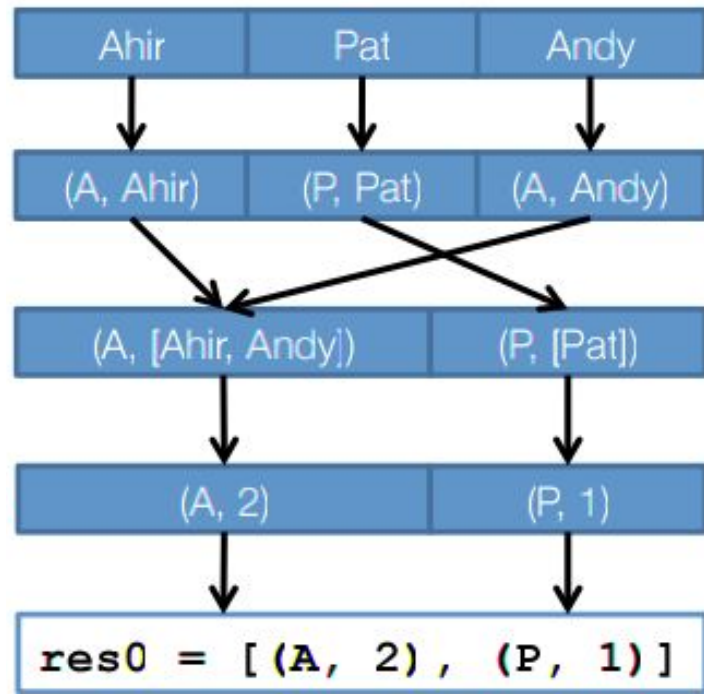
- Multiple stages - not just Map and then Reduce
 - Example: Map (stage 0), Reduce (stage 1), Map (stage 2)
- Memory cache
- Container (JVM) reuse
- Intermediate data contains serialized RDDs on disk

Yet another coding example

Another pass at computing in Spark

Find number of distinct names per “first letter”

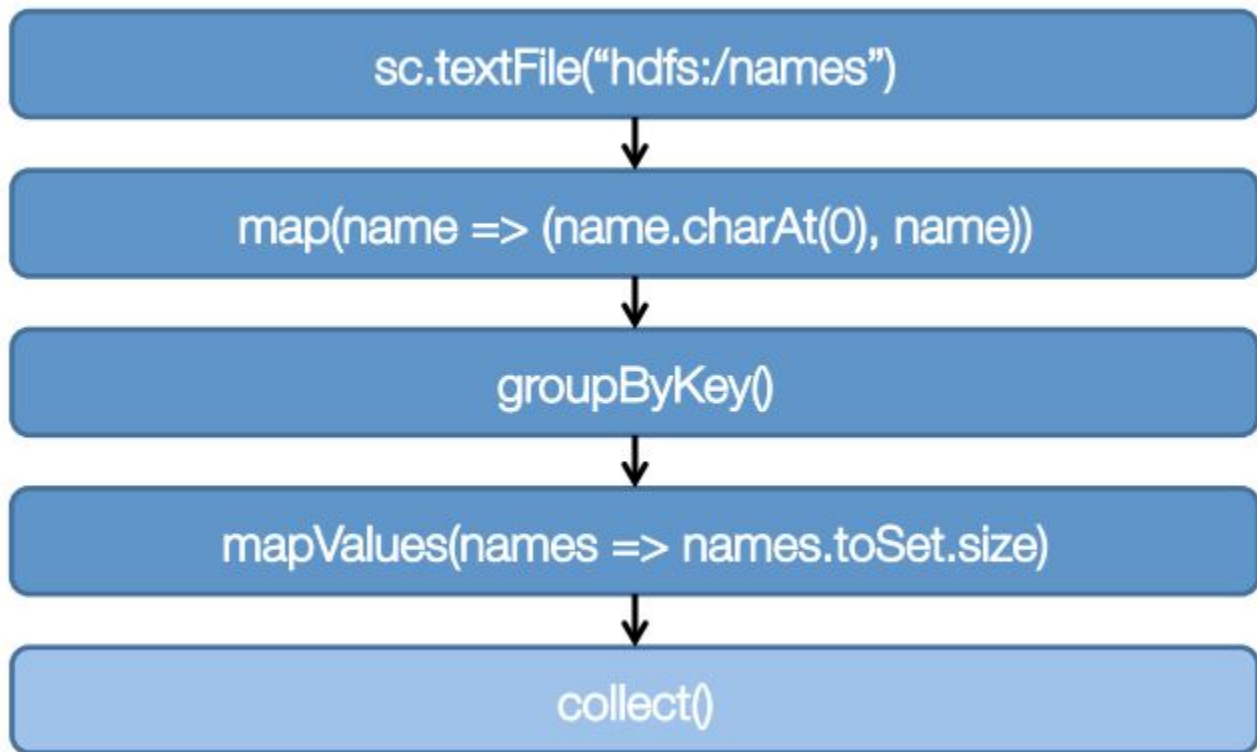
```
sc.textFile("hdfs:/names")  
  .map(name=>(name.charAt(0), name))  
  .groupByKey()  
  .mapValues(names=>names.toSet.size)  
  .collect()
```



Spark Execution Model

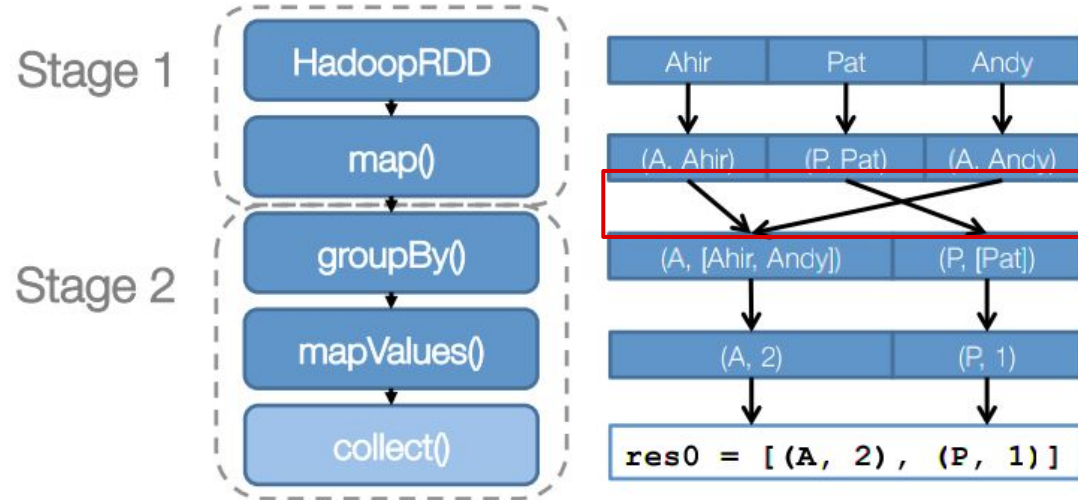
1. Create DAG of RDDs to represent computation
2. Create logical execution plan for DAG
3. Schedule and execute individual tasks

Step 1: DAG for job



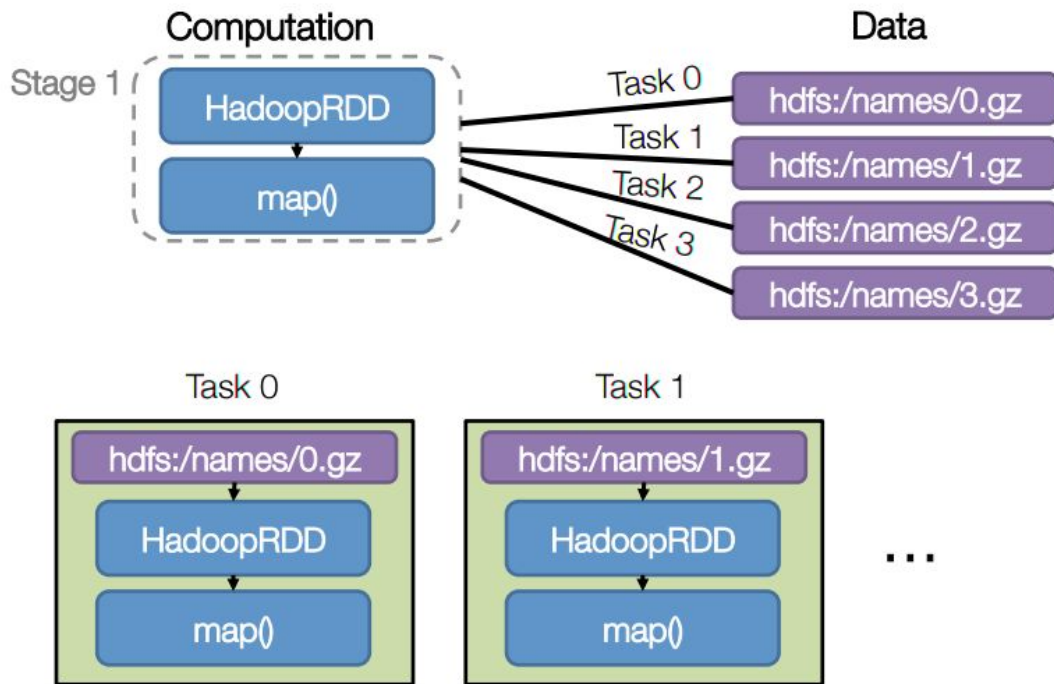
Step 2: Execution plan

- Pipeline as much as possible
- Split into “**stages**” based on need to reorganize data




Step 3: Scheduling

- Split stages into tasks
- A task is data + computation
- Execute each task in a stage before moving on



Shuffle

1. Bucket up the data: Hash by key into buckets
2. Write buckets to disk
3. Pull bucket files to nodes used for stage 2



Optimizations

Problem: Ran the `toSet()` operation, to remove duplicates, at end

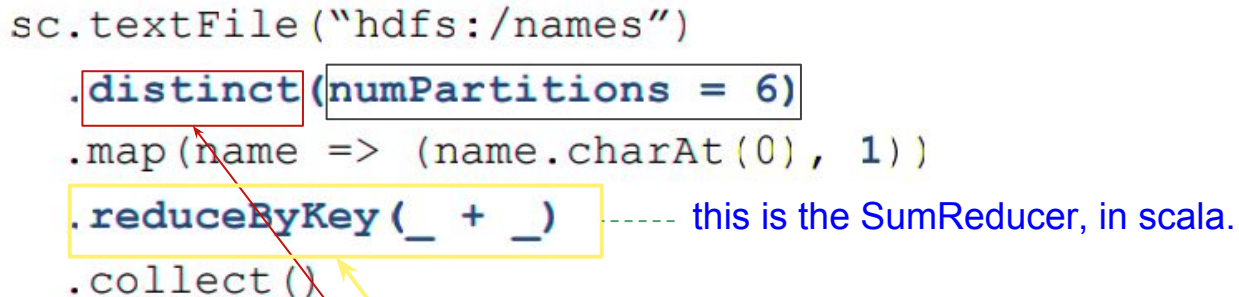
- probably want to cull the data earlier

Problem: May not be enough concurrency

- Need “reasonable number” of data partitions
- Commonly between 100 and 10,000 partitions
- Lower bound: At least $\sim 2x$ number of cores in cluster
- Upper bound: Ensure tasks take at least 100ms

Revised code with optimizations

```
sc.textFile("hdfs:/names")  
  .distinct(numPartitions = 6)  
  .map(name => (name.charAt(0), 1))  
  .reduceByKey(_ + _) ----- this is the SumReducer, in scala.  
  .collect()
```



Original:

```
sc.textFile("hdfs:/names")  
  .map(name => (name.charAt(0), name))  
  .groupByKey()  
  .mapValues { names => names.toSet.size }  
  .collect()
```

Using Mapper setup and cleanup methods is a story for another day...



Hadoop Streaming

The MaxTemp MapReducer written in Python

Special Practice

On your VM's Desktop, there is a folder called
“[hadoop-streaming](#)”.

This contains a SpecialPractice to help you learn

Hadoop Streaming

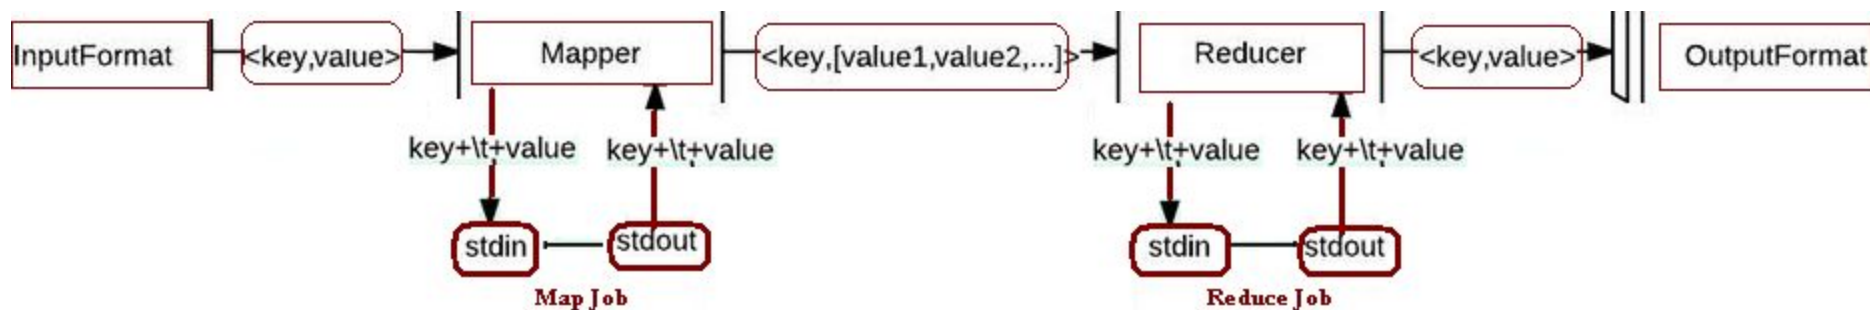
- **Features**
- **Example using Python**
 - **python mapper**
 - **python reducer**
- **How to run**

Hadoop Streaming: features

- **Run MapReduce using *any* language that can read from standard input and write to standard output.**
- **An important difference:**
 - Hadoop MapReduce functions process one record at a time
 - Hadoop Streaming functions read from stdin and control the read process.

How it works:

Streaming calls code from Mapper or Reducer



hadoop streaming: Python mapper

```
import re  
import sys
```

```
for line in sys.stdin:  
    val = line.strip()  
    (year,temp,q)=val[15:19], val[87:92], val[92:93]  
    if (temp != "+9999" and re.match("[01459]", q)):  
        print "%s\t%s" % (year, temp)
```

hadoop streaming: Python Reducer

```
import sys
```

```
(last_key, max_val) = (None, -sys.maxint)
```

```
for line in sys.stdin:
```

```
    (key,val) = line.strip().split("\t")
```

```
    if last_key and last_key != key:
```

```
        print "%s \t %s" % (last_key, max_val)
```

```
        (last_key, max_val) = (key, int(val))
```

```
    else:
```

```
        (last_key, max_val) = (key, max(max_val, int(val)))
```


```
if last_key:
```

```
    print "%s \t%s" % (last_key, max_val)
```

hadoop streaming: running the job

```
$ hadoop jar /usr/lib/hadoop-<version>-mapreduce/\
contrib/streaming/hadoop-streaming-<version>.jar \
-input inputDir -output outputDir \
-file pathToMapScript -file pathToReduceScript \
-mapper mapBasename -reducer reduceBasename
```

Hadoop supplies the jar
for streaming



Example: running hadoop streaming with Python in the studentVM

```
hadoop jar /usr/lib/hadoop-0.20-mapreduce/\
contrib/streaming/hadoop-streaming-2.0.0-mr1-cdh4.2.1.jar \
-input shakespeare -output avgwordstreaming \
-file mapper.py \
-file reducer.py \
-mapper mapper.py -reducer reducer.py
```


Key Points

- **To write a Mapper and a Reducer**
 - **can use any language that reads and writes to stdio**
 - **code must iterate through input data**
- **To run with “hadoop jar”:**
 - **use the hadoop-*-streaming.jar**
 - **use the -mapper and -reducer flags**

Extra slides

job launch details

