

Intro to Hadoop SerDe

Formats, compression and serialization/deserialization

What I learned about SerDe* from backpacking

Food -> freeze-dried -> packs well, is small and light -> rehydrated, still recognizable

analogous to

Data -> serialized and compressed -> packs well, is small -> decompressed and deserialized, still recognizable

Being able to recognize your data is important

*SerDe stands for serialization/deserialization

How do we “freeze-dry” and “rehydrate” files?

Formats

Text
SequenceFile
Avro
Parquet
ORC

Compression

Snappy
LZ4
LZF
LZO
ZLIB (Deflate)
BZip2 (splittable)

Agenda

- Native support
- Compressible
- Splittable
- Self-describing

	Avro	SequenceFile	ORC	Parquet
Origin	Hadoop	Hadoop	Hortonworks	Twitter Clouder
Rationale	Portable replacement for Writables	Need splittable, compressible input to MR2	Improved RCFile	Similar to ORC but portable
Applicability	General	General	Hive	General
Storage	Row	Row	Columnar	Columnar
Schema evolution	Yes	No	Yes	Yes
Block size (MB)	128	128	256	128/256**

Serialization/Deserialization in different applications

- Serializable types in Spark
 - **java.io.Serializable** objects
 - faster serialization with **Kryo**
 - requires **Kryo** *registration*
- Serializable types in MapReduce 2:
 - **hadoop.io.Writables**
 - **apache.Avro**
- Serialization classes in Hive
 - Parquet - ParquetSerDe
 - Avro - AvroSerDe
 - ORC - ORCSerDe
 - Text
 - JSON format - JsonSerDe
 - “Regular” format - RegExSerDe
 - Table format - OpenCSVSerDe
 - Thrift - ThriftSerDe

Serialization in Spark with an example

- averaging with aggregate requires us to keep
 - a running total of input values
 - a count of input values
- working with aggregation, simple or complex, can be simpler if we have objects that contain the information and operations needed
 - simplifies the lambda functions
 - clarifies the operations
- we can create a data class for aggregation, but it must be serializable
- if we create a data class, we should register it with Kryo

Serializable for aggregate when averaging

```
public class AvgSer implements java.io.Serializable {  
  
    public AvgSer() {  
        total_ = 0;  
        num_ = 0;  
    }  
  
    public AvgSer(float total, int num) {  
        total_ = total;  
        num_ = num;  
    }  
  
    public float avg() {  
        return total_ / (float) num_;  
    }  
  
    public float total_;  
    public int num_;  
}
```

- averaging requires us to keep
 - a running total of input values
 - a count of input values
- serializable data encapsulate info and ops
 - simplify lambda functions
 - clarify the ops

```

public final class AvgWithAvgSerAndKryo {

    public static class AvgRegistrator implements KryoRegistrator {

        @Override
        public void registerClasses(Kryo kryo) {
            kryo.register(AvgSer.class, new FieldSerializer(kryo, AvgSer.class));
        }
    }

    public static void main(String args[]) {

        SparkConf conf = new SparkConf().set("spark.serializer", "org.apache.spark.serializer.KryoSerializer");
        conf.setMaster("local");
        conf.setAppName("Avg with serialization and kryo");
        JavaSparkContext sc = new JavaSparkContext(conf);

        JavaRDD<Float> input = sc.parallelize(Arrays.asList(1f, 2f, 3f, 4f));
        AvgSer avgs = computeAvg(input);
        System.out.println("Average for list: " + avgs.avg());

        sc.stop();
    }

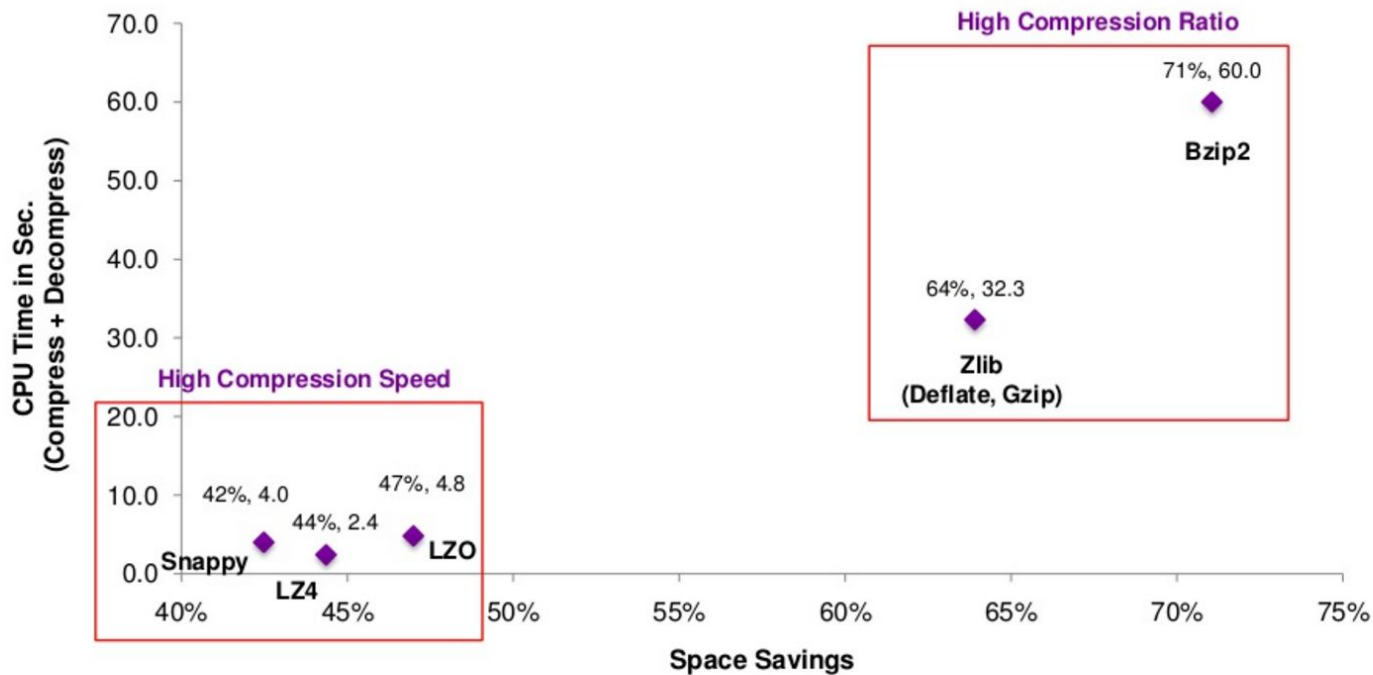
    static AvgSer computeAvg(JavaRDD<Float> input) {

        AvgSer zeroValue = new AvgSer();
        return input.aggregate(zeroValue, (a, x) -> new AvgSer(a.total_ + x, a.count_ + 1),
            (a, b) -> new AvgSer(a.total_ + b.total_, a.count_ + b.count_));
    }
}

```


Space-Time Tradeoff of Compression Options

Codec Performance on the Wikipedia Text Corpus




Other performance comparisons: [Blog comparing snappy, LZF, LZO and ZLIB](#)

Note:

A 265 MB corpus from Wikipedia was used for the performance comparisons.

Space savings is defined as $[1 - (\text{Compressed} / \text{Uncompressed})]$

Space-Time Tradeoff of Compressors



Note:
A 265 MB corpus from Wikipedia was used for the performance comparisons.
Space savings is defined as [1 - (Compressed/ Uncompressed)]

Algorithm	Compression Ratio	Speed	Splittable
Snappy	~40%	4.0	NO
LZ4	~45%	2.4	NO
LZF	~45%	NA	NO
LZO	~45%	4.8	NO
ZLIB (Deflate)	~60%	32.8	NO
BZip2	~71%	60.0	YES

```

public class CompressingAndSavingFiles {

    public static void main(String[] args) throws Exception {

        if (args.length != 1) {
            System.out.println("For running in Eclipse - the argument is:  outputFile");
            System.exit(-1);
        }

        String output1 = args[0] + "deflated";
        String output2 = args[0] + "snappy";

        JavaSparkContext sc = new JavaSparkContext("local", "saving compressed data to file");

        JavaRDD<Integer> rdd = sc.parallelize(Arrays.asList(1, 2, 3, 4, 5, 6, 7));

        JavaRDD<Integer> result = rdd.map(x -> x * x).filter(x -> x % 2 != 0);

        System.out.println("lambda results:  " + result.collect());

        /*-
        * choices for codecs descending from org.apache.hadoop.io.compress.CompressionCodec
        *
        * DeflateCodec (resolves to GZipCodec)
        *
        * Lz4Code
        *
        * SnappyCodec
        *
        * BZip2Codec
        */
        Class codec1 = org.apache.hadoop.io.compress.DeflateCodec.class;
        result.saveAsTextFile(output1, codec1);

        Class codec2 = org.apache.hadoop.io.compress.SnappyCodec.class;
        result.saveAsTextFile(output1, codec2);

        sc.stop();
    }
}

```

Hey, wait, what about...

..writing files using

- Spark RDDs, DataFrames, DataSets
- Hadoop IO standalone

What about the wrappers for Hadoop IO in Spark?

How do we DO anything with the serialization framework...

Ah, the glories of transitions and evolving open-source.

Reference slides

		Writables	Thrift	Protocol Buffers	Avro	RCFile	ORC	Parquet
Origin		Hadoop	Facebook	Google	Hadoop	Facebook et al.	Hortonworks	Twitter Cloudera
Rationale			Language-independent interfaces (IDL) to services	Data exchange between services through IDL	Portable replacement for Writables	Columnar replacement for SequenceFiles	Improved RCFile	Similar to ORC but portable
Applicability		Java	General	General	General	Hive	Hive	General
Storage		Row	Row	Row	Row	Columnar	Columnar	Columnar
Native MR support		Yes	No	No	Yes	Yes	Yes	Yes
Compressible		Yes	No	No	Yes	Yes	Yes	Yes
Splittable		Yes	No	No	Yes	Yes	Yes	Yes
Schema evolution		No	Yes	Yes	Yes	No	No	Yes**
Self-describing		No	No	No	Yes	Yes	Yes	Yes
Block size (MB)		-	-	-	64	4	256	128/256***
Supported data types	Simple	Null	✓	✓	✓	✓	✓	✓
		Boolean	✓	✓	✓	✓	✓	✓
		Integer (8 bits): byte	✓	✓	✓	✓	✓	✓*
	Numerical	Integer (16 bits): short	✓	✓	✓*	✓	✓	✓*
		Integer (32 bits): int	✓	✓	✓	✓	✓	✓*
		Integer (64 bits): long	✓	✓	✓	✓	✓	✓*
		Float (32 bits): float	✓	✓	✓	✓	✓	✓
		Float (64 bits): double	✓	✓	✓	✓	✓	✓
	Text	String	✓	✓	✓	✓	✓	✓
		Character				✓	✓	
		Date				✓	✓	✓
	Time	Timestamp				✓	✓	✓
		Interval						✓
	Binary		✓					✓
		Array	✓		✓	✓	✓	
		List		✓			✓	✓
	Collections	Set	✓					
		Map	✓	✓	✓	✓	✓	✓
		SortedMap	✓					
		Enum	✓	✓	✓			
		Tuple	✓					
		Union			✓	✓	✓	
	Records	Fixed			✓			
		Struct		✓		✓	✓	
		Object/Record	✓	✓	✓			

* Including unsigned data types / ** Append only / *** 256 MB in Impala

Compression types for different file formats

FileType	Compression
Text	LZO
Avro	Snappy, GZIP, deflate, BZIP2
Parquet	Snappy, GZIP, deflate, BZIP2
ORC	Snappy, ZLib
RCFile	Snappy, GZIP, deflate, BZIP2
SequenceFile	Snappy, GZIP, deflate, BZIP2

		Avro	RCFile	ORC	Parquet
Origin		Hadoop	Facebook et al.	Hortonworks	Twitter Cloudera
Rationale		Portable replacement for Writables	Columnar replacement for SequenceFiles	Improved RCFile	Similar to ORC but portable
Applicability		General	Hive	Hive	General
Storage		Row	Columnar	Columnar	Columnar
Schema evolution		Yes	No	No	Yes**
Self-describing		Yes	Yes	Yes	Yes
Block size (MB)		64	4	256	128/256***
Supported data types	Simple	Null	✓	✓	✓
		Boolean	✓	✓	✓
	Numerical	Integer (8 bits): byte	✓	✓	✓*
		Integer (16 bits): short		✓	✓*
		Integer (32 bits): int	✓	✓	✓*
		Integer (64 bits): long	✓	✓	✓*
		Float (32 bits): float	✓	✓	✓
	Text	Float (64 bits): double	✓	✓	✓
		String	✓	✓	✓
		Character		✓	
	Time	Date		✓	✓
		Timestamp		✓	✓
		Interval			✓
	Binary				✓
		Array	✓	✓	
	Collections	List		✓	✓
		Set			
		Map	✓	✓	✓
		SortedMap			
		Enum	✓		
		Tuple			
		Union	✓	✓	
	Records	Fixed	✓		
		Struct		✓	
		Object/Record	✓		

Each format supports
different data types

Hadoop IO - loading data with input formats in Spark

To use TextInputFormat:

```
lines = sc.textFile(pathStr);
```

To use SequenceFileInputFormat:

```
records = sc.sequenceFile(pathStr, key.class, value.class);
```

EXAMPLE: READING A SEQUENCE FILE

```
import org.apache.hadoop.io.IntWritable
```

```
import org.apache.hadoop.io.Text
```

```
JavaPairRDD rdd = sc.sequenceFile("mySequenceFile", IntWritable.class,  
Text.class)
```

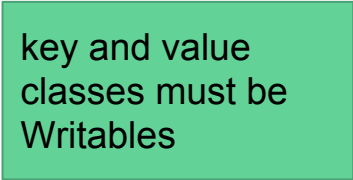
Hadoop IO - saving in different formats in spark

To use TextInputFormat:

```
sc.saveAsTextFile(rdd);
```

To use SequenceFileInputFormat:

```
sc.saveAsHadoopFile(fileName,  
    <key>.class,  
    <value>.class,  
    SequenceFileOutputFormat.class);
```

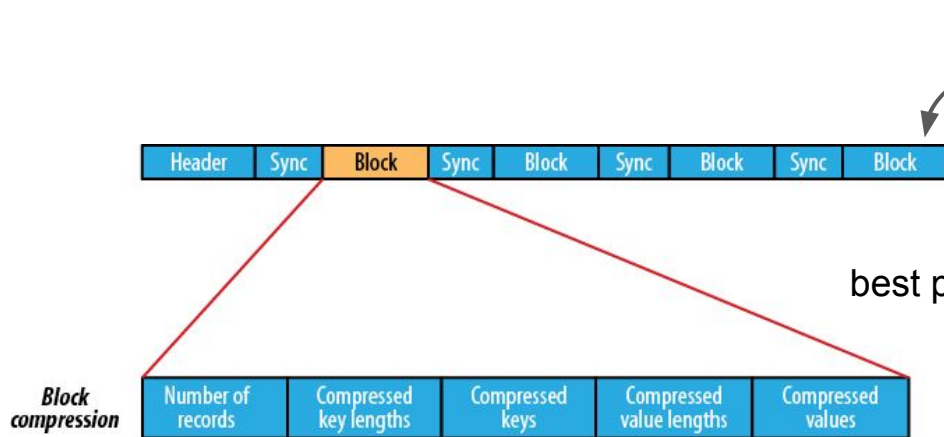


key and value
classes must be
Writable

Sequence files

- Store key, value pairs as records
- Compressible, splittable
- Self-aware header

writer.append(key, value);




best practice: create with block compression

Examples of reading and writing SequenceFiles

See class-examples-io

- **SaveAsSequenceFile.java** - run this first, it will create a sequence file
 - To run, in Run Configurations, specify the name of the file. For example:
data/panda.seq
 -
- **LoadSequenceFile.java** - run this next, it will load the sequence file and show you the contents

using wholeTextFiles: Handling small files



- `sc.wholeTextFiles(directory)`
 - Maps entire contents of each file in directory to a single RDD element
 - Works only for small files (element must fit in memory)



```
(file1.json, {"firstName": "Fred", "lastName": "Flintstone", "userid": "123"})  
(file2.json, {"firstName": "Barney", "lastName": "Rubble", "userid": "234"})  
(file3.xml, ... )  
(file4.xml, ... )
```

```
public class ReadWholeTextFiles {  
    public static void main(String[] args) {  
        if (args.length != 1) {  
            System.out.println("For running in Eclipse - the argument is: inputDir");  
            System.exit(-1);  
        }  
  
        String input = args[0];  
  
        JavaSparkContext sc = new JavaSparkContext("local", "read whole text files");  
        JavaPairRDD<String, String> rdd1 = sc.wholeTextFiles(input);  
        Map<String, String> map = rdd1.collectAsMap();  
  
        Set<String> keys = map.keySet();  
  
        for (String key : keys) {  
            System.out.println("Key: " + key);  
  
            /*-  
             * include to see all the data :  
             * System.out.println("File contents: " + map.get(key));  
             */  
        }  
    }  
}
```

Examples of reading WholeTextFiles

See [class-examples-io](#)

- [ReadWholeTextFiles.java](#) - run this first with a directory of small files
 - To run, in Run Configurations, specify the name of the directory.
 - For example: **data/stock-sample-tiny**

This will read each file in the directory and show you the files it has read.

Hadoop / Spark Input formats (1)

Avro files

```
newAPIHadoopFile(path, InputFormat class, key.class, value.class, [conf])
```

```
JavaPairRDD avroRecords = sc.newAPIHadoopFile(input,  
AvroKeyInputFormat.class, Text.class, NullWritable.class,  
new Configuration());
```


Hadoop /Spark Input formats (2)

KeyValueText files

```
newAPIHadoopFile(path, InputFormat class, key.class, value.class, [conf])
```

```
Configuration conf = new Configuration();  
conf.set("mapreduce.input.keyvaluelinerecordreader.key.value.separator",  
        ":");  
JavaPairRDD keyValues = sc.newAPIHadoopFile(input,  
        KeyValueTextInputFormat.class, Text.class, IntWritable.class,  
        conf)
```

Hadoop / Spark input and output formats

OK, I'm in the 3rd circle of Hell...

I hope I'm not alone.

Avro data

- Useful for evolving binary data, JSON based.
 - Schema evolution – Avro requires schemas written in JSON. Can use different schemas for serialization and deserialization.
 - Untagged data – Allows each datum be written without overhead: more compact data encoding, and faster data processing.
 - Dynamic typing – serialization and deserialization without code generation.

Avro file - companyInfo.avsc

```
{
  "type": "record",
  "name": "CompanyInfo",
  "fields": [
    {"name": "Symbol", "type": "string"},
    {"name": "Name", "type": "string"},
    {"name": "LastSale", "type": "string"},
    {"name": "MarketCap", "type": "string"},
    {"name": "IPOyear", "type": "string"},
    {"name": "Sector", "type": "string"},
    {"name": "Industry", "type": "string"},
    {"name": "SummaryQuoteURL", "type": "string"}
  ]
}
```

This file is in src/main/resources/avro

```

private static void readAvro(File inputFile, int howMany) throws IOException {
    System.out.println("***** readAvro from " + inputFile + " for " + howMany + " *****");
    DatumReader<GenericRecord> reader = new GenericDatumReader<GenericRecord>();
    DataFileReader<GenericRecord> dataFileReader = new DataFileReader<GenericRecord>(inputFile, reader);

    Schema companyInfoSchema = dataFileReader.getSchema();
    System.out.println("CompanyInfo schema from data file: " + companyInfoSchema.toString(true));

    GenericData.Record record = new GenericData.Record(companyInfoSchema);
    int counter = 0;
    while (dataFileReader.hasNext()) {
        dataFileReader.next(record);
        System.out.println(record.get("Symbol").toString() + "\t" + record.get("Name").toString() + "\t"
            + record.get("LastSale").toString() + "\t" + record.get("MarketCap").toString() + "\t"
            + record.get("IPOyear").toString() + "\t" + record.get("Sector").toString() + "\t"
            + record.get("Industry").toString() + "\t" + record.get("SummaryQuoteURL").toString());

        counter++;

        if (counter == howMany) {
            break;
        }
    }
    dataFileReader.close();
}

```

Create reader
for Avro file

Create record

Write records to
System.out

Read and parse
schema (.avsc)

Create writer

Create record

Read text fields
and write to
record

```
private static void writeAvro(File inputFile, File outputFile) throws IOException {  
    InputStream schemaIS = CompanyInfoAvroUtility.class  
        .getResourceAsStream(SchemaConstants.COMPANYINFO_AVRO_SCHEMA);  
    if (schemaIS == null) {  
        throw new IllegalStateException("Unable to find " + SchemaConstants.COMPANYINFO_AVRO_SCHEMA);  
    }  
    Schema companyInfoSchema = new Parser().parse(schemaIS);  
  
    DatumWriter<GenericRecord> writer = new GenericDatumWriter<GenericRecord>(companyInfoSchema);  
    DataFileWriter<GenericRecord> dataFileWriter = new DataFileWriter<GenericRecord>(writer);  
    dataFileWriter.setCodec(CodecFactory.deflateCodec(9));  
    dataFileWriter.create(companyInfoSchema, outputFile);  
  
    BufferedReader reader = new BufferedReader(new FileReader(inputFile));  
    String line = null;  
    GenericData.Record record = new GenericData.Record(companyInfoSchema);  
  
    System.out.println("***** writing data in Avro format *****");  
    while ((line = reader.readLine()) != null) {  
        String[] tokens = line.split(regex1, -1);  
  
        for (int i = 0; i < tokens.length; i++)  
            tokens[i] = tokens[i].replace("\\", "");  
  
        if (tokens.length == 9) {  
            record.put("Symbol", new Utf8(tokens[0]));  
            record.put("Name", new Utf8(tokens[1]));  
            record.put("LastSale", new Utf8(tokens[2]));  
            record.put("MarketCap", new Utf8(tokens[3]));  
            record.put("IPOyear", new Utf8(tokens[4]));  
            record.put("Sector", new Utf8(tokens[5]));  
            record.put("Industry", new Utf8(tokens[6]));  
            record.put("SummaryQuoteURL", new Utf8(tokens[7]));  
            dataFileWriter.append(record);  
        }  
    }  
  
    reader.close();  
    dataFileWriter.close();  
  
    System.out.println("***** finished writing data in Avro format to " + outputFile + " *****");  
}
```

To run example:

click on `CompanyInfoAvroUtility.java`

select “Run As” -> select “Run Configurations”

Arguments:

- To convert files:

```
convert data/companylistNASDAQ.csv data/companylistNASDAQ.avro
```

- To read 25 files:

```
read data/companylistNASDAQ.avro 25
```

For a deeper (but possibly out-of-date??) analysis of data flows in Hadoop

Hadoop Application Architectures

Designing Real-World Big Data Applications

By [Mark Grover](#), [Ted Malaska](#), [Jonathan Seidman](#), [Gwen Shapira](#)

Publisher: O'Reilly Media

Final Release Date: June 2015

Pages: 400

[O'Reilly link - to buy or browse](#)

Still to discuss: Parquet

Serialization in Hadoop

Hadoop IO can serialize and deserialize app data of the following types:

- Java Serializable, with Kryo (Spark only)
- Avro
- Parquet
- ORC
- Writable
- SequenceFile

Used by Hadoop (internally) but not by apps in Spark, Hive or MR2

- Thrift
- Protocol buffers