## HDFS Cheat Sheet by Matthew Rathbone

Commonly used commands for HDFS. This is not an exhaustive list, but the most useful ones on a day to day basis.

Listing Files		
This is how you inspect HDFS to see what it contains. Use these commands to find files and their associated metadata.		
hadoop fs -ls ./example	List files in a directory	
hadoop fs -ls	List files matching a pattern	
./example/matthew*		
hadoop fs -ls -h ./example/	List files with human-friendly file sizes (eg 1.3M vs 1331325)	
hadoop fs -ls -R ./example/	Recursively list files in this directory AND all child directories.	
hadoop fs -ls -d ./example/	List files, but with directories shown as files. So in this case it will show info about the example folder itself.	

Uploading/Downloading Files	
Moving files between HDFS and t	he local filesystem and back.
hadoop fs -put ./localfile.txt ./example/	Upload a file from your local machne to a specific directory on HDFS.
hadoop fs -put -f ./localfile.txt ./example/	Upload a file and overwrite any existing file on HDFS.
hadoop fs -put -l ./localfile.txt ./example/	Upload a file and set a replication factor of 1 (you probably shouldn't ever really use this).
hadoop fs -get ./example/remotefile.txt ./	Download a file from HDFS to your local machine.
hadoop fs -get -p ./example/remotefile.txt ./	Download a file from HDFS to your local machine, preserving metadata (eg modified time).
hadoop fs -get ./example/*.txt ./	Download a set of files that match a pattern to your local machine.

Reading & Writing Files	
Reading file contents without do	wnloading the file itself.
hadoop fs -text ./example/file.txt	Print the contents of a file to the terminal, decompressing if nessecery.
hadoop fs -cat ./example/*.txt	Print to the terminal the contents of all files that match the provided pattern. Note - this will NOT decompress like 'text' will.
hadoop fs [-cat,-text] - ignoreCrc ./example/*.txt	As above, but disable the verification checksum.
hadoop fs -appendToFile ./localfile.txt ./example/remotefile.txt	Append the contents of a local file to a file on HDFS. This is only supported in Hadoop versions 2.1.1+

File Management  Organize your files with these commands.		
./example/f2.txt	menante to name it the same)	
hafoop fs -cp	copy a file to a different file/directory (omit the	
./example/f1.txt	filename to name it the same)	
./example/f2.txt		
hafoop fs -rm	Delete a file (sends it to the trash)	
./example/f1.txt		
hafoop fs -rm -skipTrash	Actually delete the file. No trash.	
./example/f1.txt		
hafoop fs -rm -r	Recursively delete a directory and it's contents	
./example/directory		
hafoop fs -touchz	Create a zero-length file (great for creating	
./example/somefile	_SUCCESS files).	
hafoop fs -touchz	Create a zero-length file (great for creating	
./example/somefile	_SUCCESS files).	

HDFS Administration	
Check on important stuff that is le	ess about your files and more about HDFS.
hadoop fs -df -h ./example	Show capacity and used space of the filesystem. Will show partition space remaining if you have partitions.
hadoop fs -du -h ./example/*.txt	show the amount of space used by matching files
hadoop fs -expunge	Empty the Trash (useful if you -rm without - skipTrash)
hadoop fs -chown owner:group ./example	Change ownership of a file (use -R for the directory)
hadoop fs -chmod 0700 ./example/file.txt	Change the mode of the file (eg to 0700)
hadoop fs -checksum ./example/*.txt	Fetch checksum information for the matching files (requires a datanode roundtrip, slow and intensive).